

# RASED: A Scalable Dashboard for Monitoring Road Network Updates in OSM

Mashaal Musleh, Mohamed F. Mokbel

University of Minnesota

June 2022



UNIVERSITY OF MINNESOTA  
Driven to Discover®

# Outline

- 1. Motivation**
- 2. RASED Queries**
- 3. RASED Architecture & Main Modules**
  - Data Collection & Processing
  - Storage & Indexing
  - Query Execution
  - User Interface
- 4. Experimental Results**
- 5. RASED in Action (A Quick Demo)**

# Outline

## 1. Motivation

## 2. RASED Queries

## 3. RASED Architecture & Main Modules

- Data Collection & Processing
- Storage & Indexing
- Query Execution
- User Interface

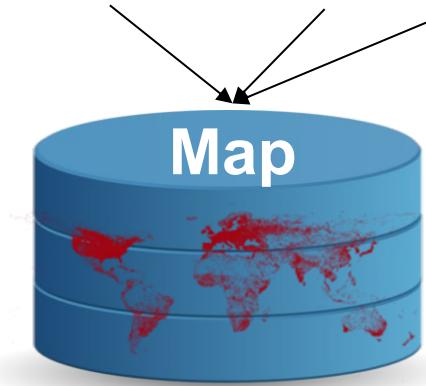
## 4. Experimental Results

## 5. RASED in Action (A Quick Demo)

# Road Network Queries...



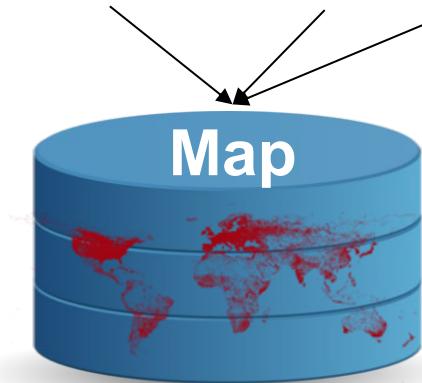
Range Query    Routing    ...



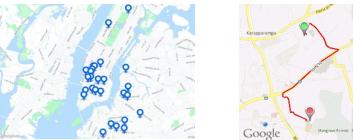
# Road Network Queries...



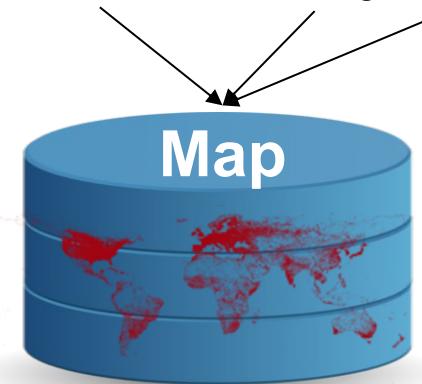
Range Query    Routing    ...



# Road Network Queries...



Range Query    Routing    ...



Map Analyzers



Quality  
Assessment

## Transactions in GIS

Full Access



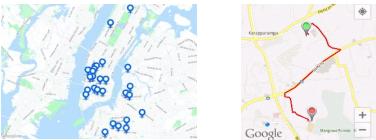
Quality Assessment of the French OpenStreetMap Dataset

Jean-François Girres , Guillaume Touya

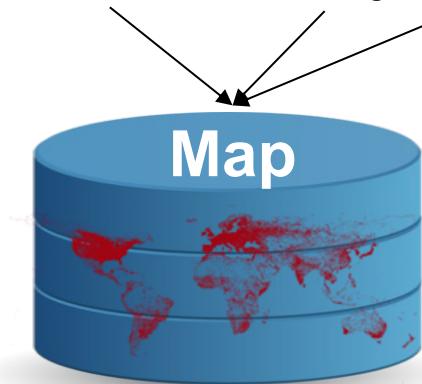
6-2018

Accuracy Evaluation of the Canadian  
OpenStreetMap Road Networks

# Road Network Queries...



Range Query    Routing    ...



Map Analyzers



Quality Assessment

- OSM:**
- 12+ Billion updates
  - 3 TB Data

## Transactions in GIS

| Full Access



Quality Assessment of the French OpenStreetMap Dataset

Jean-François Girres✉, Guillaume Touya

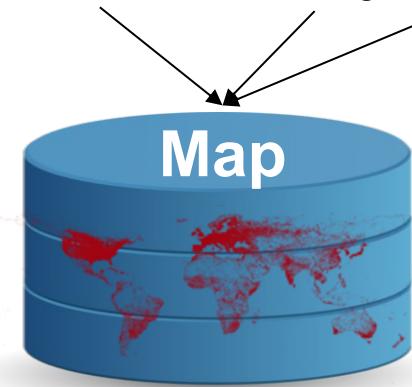
6-2018

Accuracy Evaluation of the Canadian  
OpenStreetMap Road Networks

# Road Network Queries...



Range Query      Routing      ...



Map Analyzers



Quality Assessment

- OSM:**
- 12+ Billion updates
  - 3 TB Data

**Transactions in GIS**

Full Access

**Quality Assessment of the French OpenStreetMap Dataset**  
Jean-François Girres, Guillaume Touya

6-2018

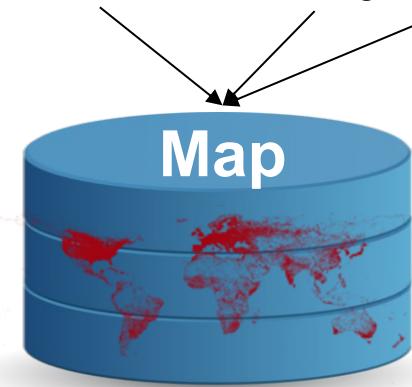
**Accuracy Evaluation of the Canadian OpenStreetMap Road Networks**

**Not Scalable**

# Road Network Queries...



Range Query      Routing      ...



## OSM:

- 12+ Billion updates
- 3 TB Data



Not Scalable

## RASED

- Global-Scale
- Comprehensive
- Highly Interactive



# How does RASED Help Map Analyzers?

# How does RASED Help Map Analyzers?

- Road changes - > gives road network stability
- Query and visualize various statistics
- Numerous query scenarios
- 15 years worth of data
- Scalable, Interactive dashboard

# How does RASED Help Map Analyzers?

- Road changes - > gives road network stability
- Query and visualize various statistics
- Numerous query scenarios
- 15 years worth of data
- Scalable, Interactive dashboard

RASED is the first-ever attempt to quantify and visualize all OSM worldwide changes on a daily basis

# Outline

- 1. Motivation**
- 2. RASED Queries**
- 3. RASED Architecture & Main Modules**
  - Data Collection & Processing
  - Storage & Indexing
  - Query Execution
  - User Interface
- 4. Experimental Results**
- 5. RASED in Action (A Quick Demo)**

# RASED Analysis Queries

Example UpdateList U

Date	Country	RoadType	ElementType	UpdateType	Latitude	Longitude	ChangesetID
06/07/2022	USA	Service	Way	New	44.5	-93.2	5582147
...	...	...	...	...			

# RASED Analysis Queries

Example UpdateList U

Date	Country	RoadType	ElementType	UpdateType	Latitude	Longitude	ChangesetID
06/07/2022	USA	Service	Way	New	44.5	-93.2	5582147
...	...	...	...	...			

- **Aggregations:** on any highlighted fields.

# RASED Analysis Queries

Example UpdateList U

Date	Country	RoadType	ElementType	UpdateType	Latitude	Longitude	ChangesetID
06/07/2022	USA	Service	Way	New	44.5	-93.2	5582147
...	...	...	...	...			

- **Aggregations:** on any highlighted fields.
- **Filters:** Temporal (e.g., time of update), spatial (e.g., country or state), road types, and update types.

# RASED Analysis Queries

Example UpdateList U

Date	Country	RoadType	ElementType	UpdateType	Latitude	Longitude	ChangesetID
06/07/2022	USA	Service	Way	New	44.5	-93.2	5582147
...	...	...	...	...			

- **Aggregations:** on any highlighted fields.
- **Filters:** Temporal (e.g., time of update), spatial (e.g., country or state), road types, and update types.
- **Goal:** find number and percentage of OSM updates:
  - E.g., per country, comparison between countries, types of updated roads, and temporal evolution of updates.

# RASED Analysis Queries

# RASED Analysis Queries

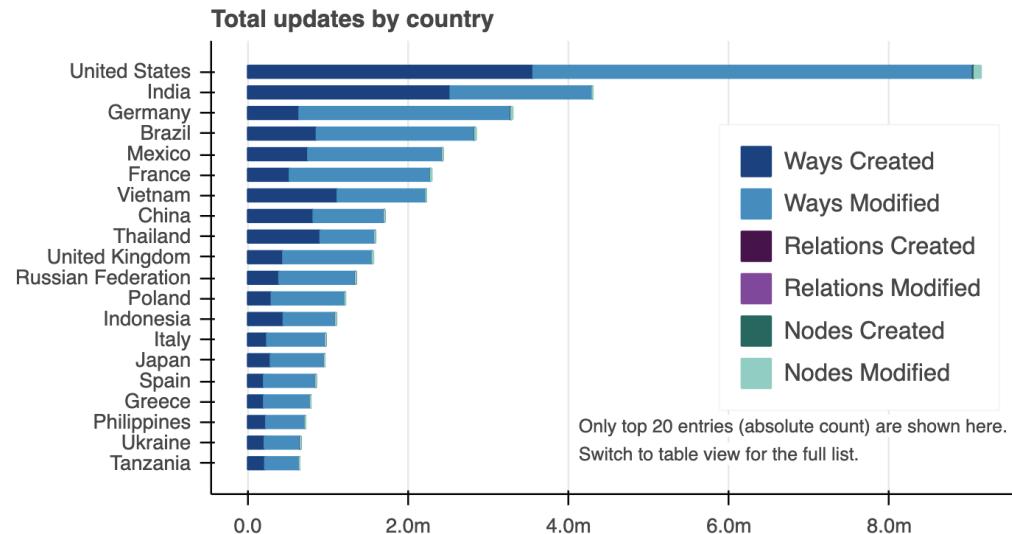
## Example 1:

Find the number of newly created or modified element types (node, way, relation) for each country road network in 2021.

# RASED Analysis Queries

## Example 1:

Find the number of newly created or modified element types (node, way, relation) for each country road network in 2021.

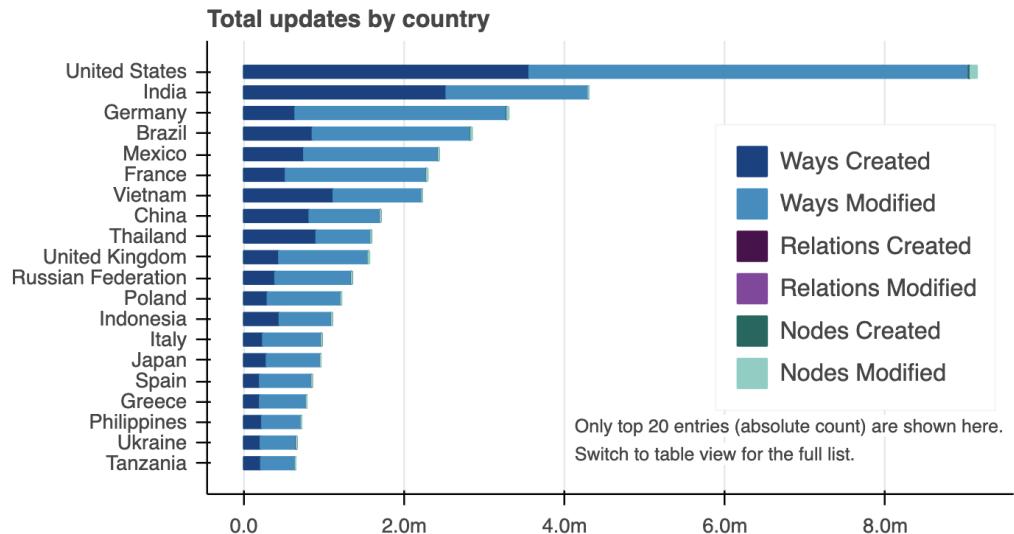


# RASED Analysis Queries

## Example 1:

Find the number of newly created or modified element types (node, way, relation) for each country road network in 2021.

```
SELECT U.Country, U.ElementType, COUNT(*)  
FROM UpdateList U  
WHERE U.Date BETWEEN 2021-01-01  
    AND 2021-12-31  
    AND U.UpdateType IN [New, Update]  
GROUP BY U.Country, U.ElementType
```

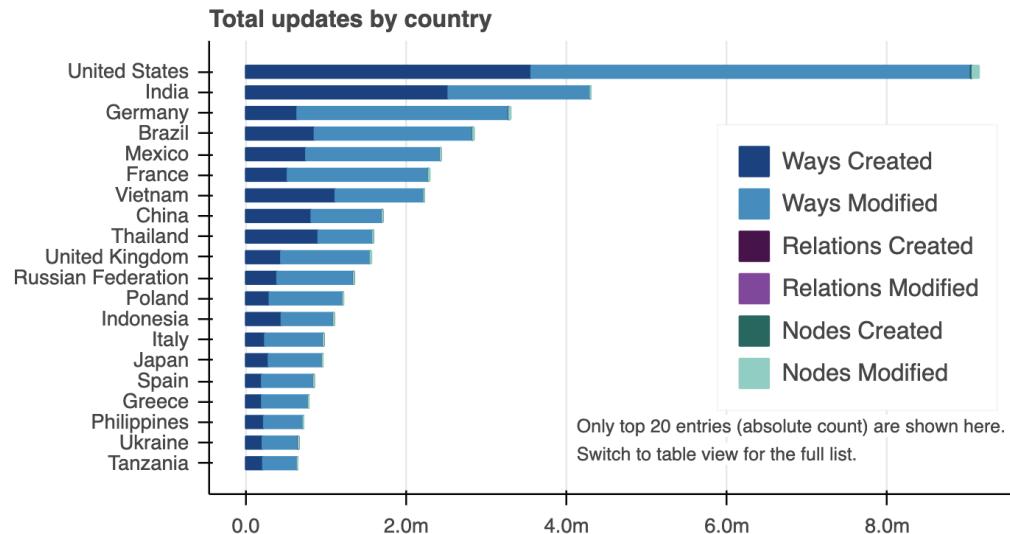


# RASED Analysis Queries

## Example 1:

Find the number of newly created or modified element types (node, way, relation) for each country road network in 2021.

```
SELECT U.Country, U.ElementType, COUNT(*)  
FROM UpdateList U  
WHERE U.Date BETWEEN 2021-01-01  
    AND 2021-12-31  
    AND U.UpdateType IN [New, Update]  
GROUP BY U.Country, U.ElementType
```



RASED queries are always supported in the order of milliseconds regardless of how large is query period.

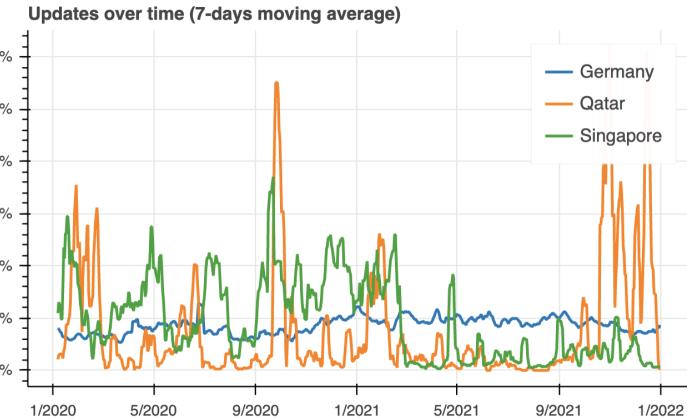
# RASED Analysis Queries



# RASED Analysis Queries

## Example 2:

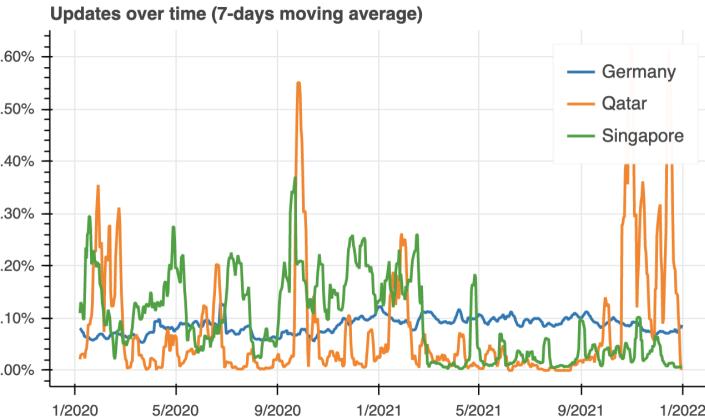
Compare the percentage of daily changes in road network in Germany, Singapore, and Qatar over 2020 and 2021.



# RASED Analysis Queries

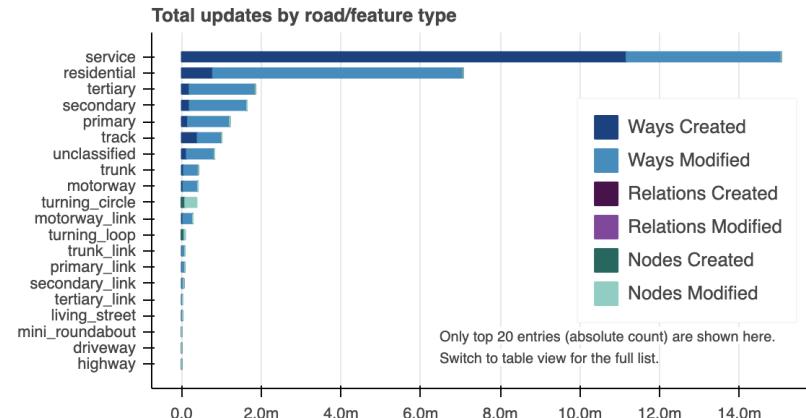
## Example 2:

Compare the percentage of daily changes in road network in Germany, Singapore, and Qatar over 2020 and 2021.



## Example 3:

Find the number of newly created or modified elements types (node, way, relation) for each road type in USA since 2018.



# Outline

**1. Motivation**

**2. RASED Queries**

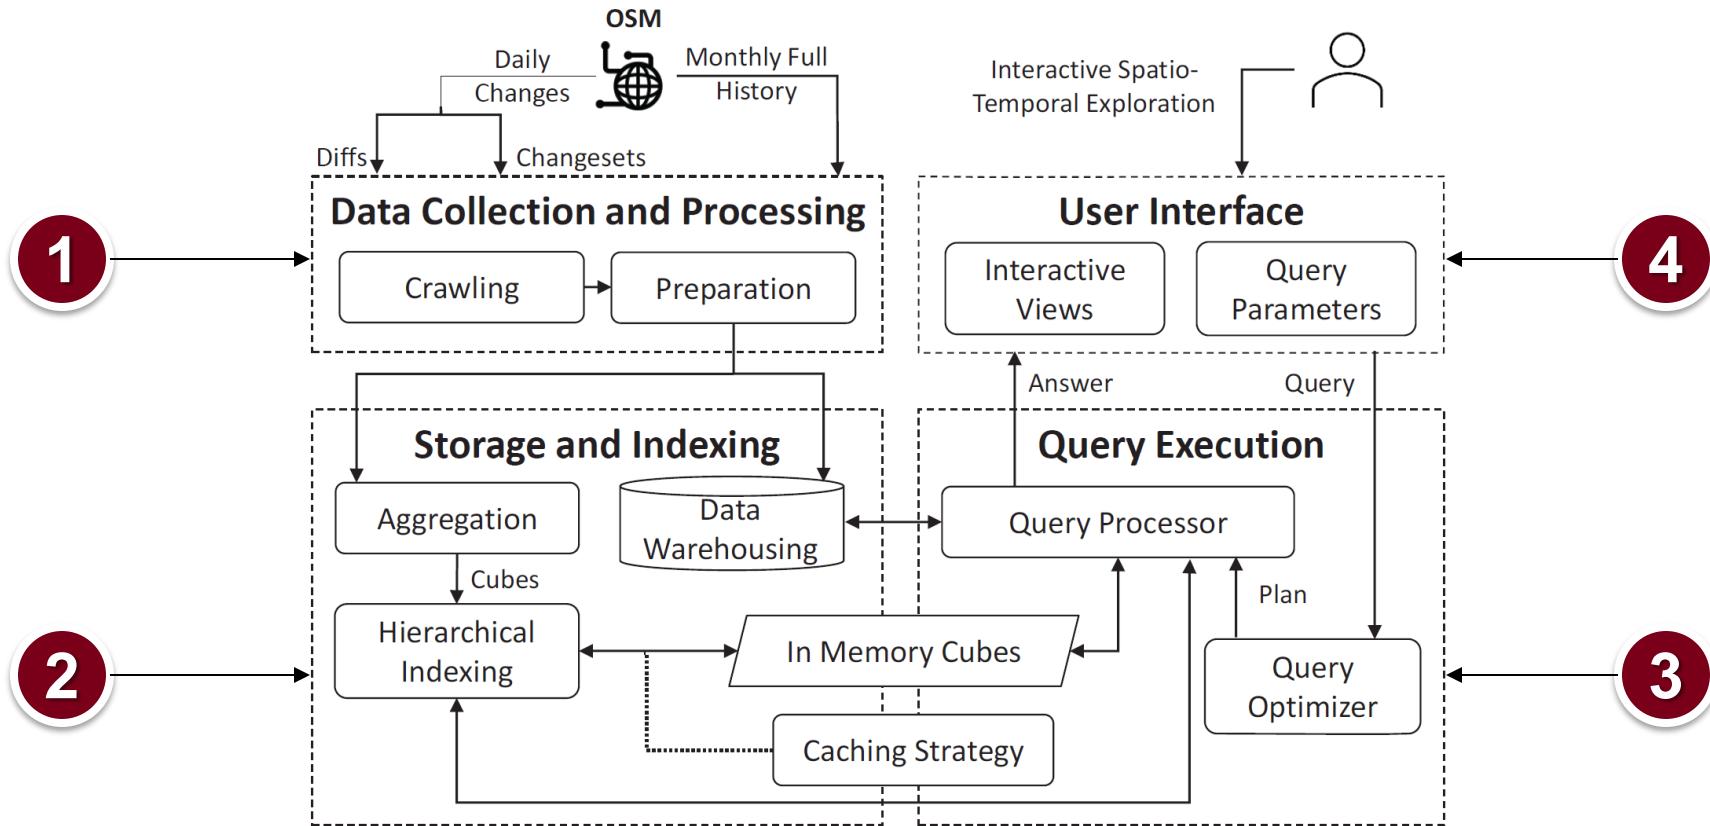
**3. RASED Architecture & Main Modules**

- Data Collection & Processing
- Storage & Indexing
- Query Execution
- User Interface

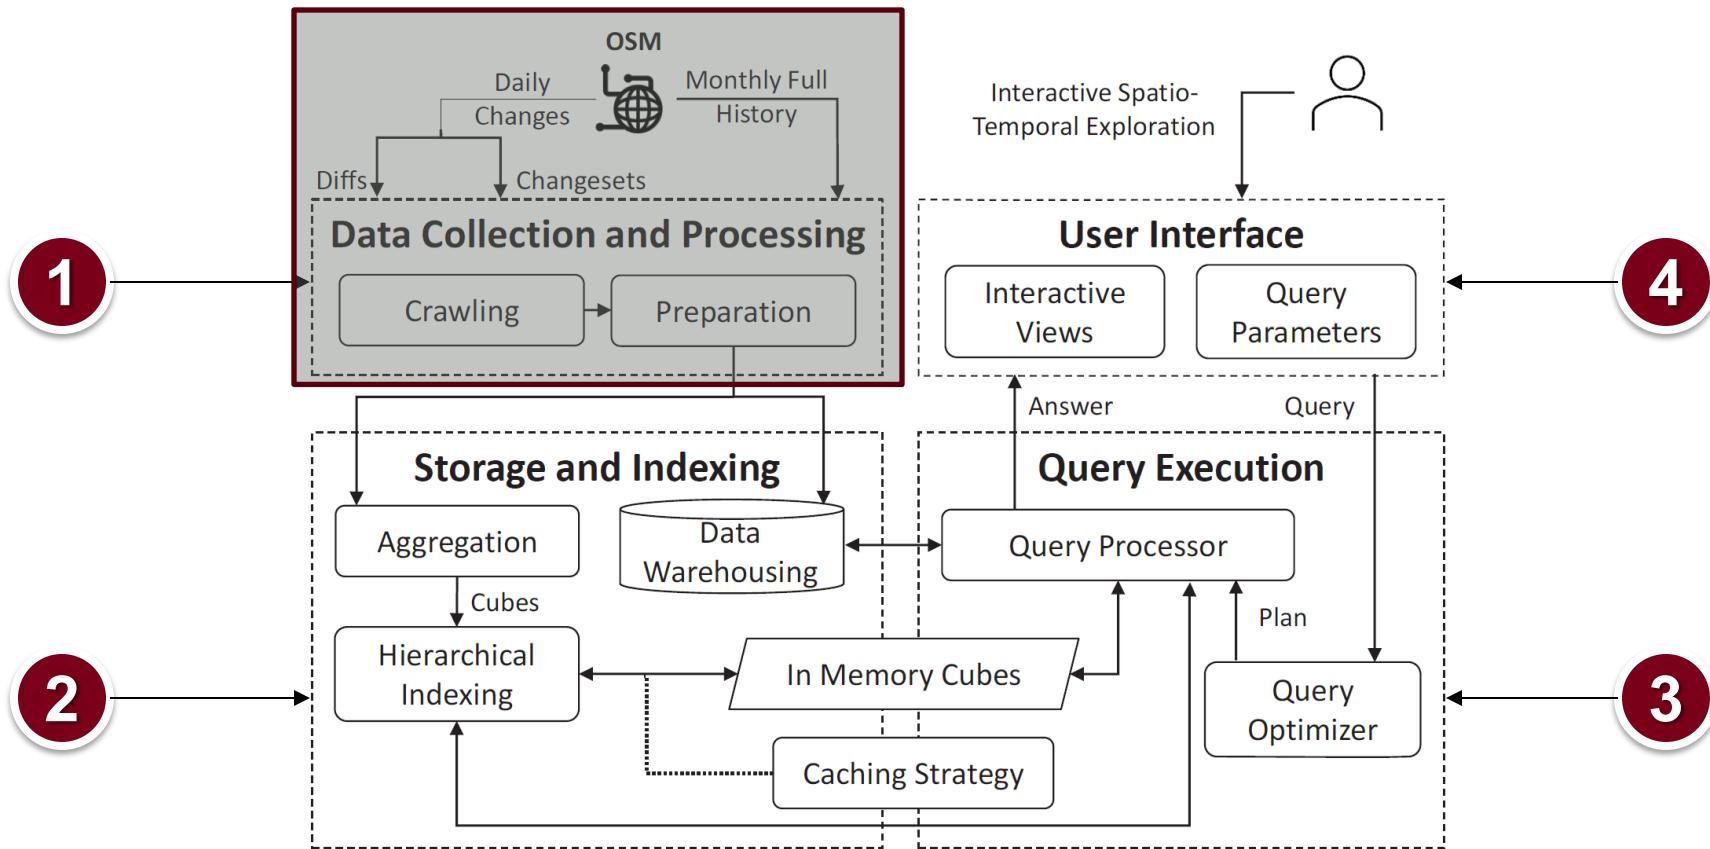
**4. Experimental Results**

**5. RASED in Action (A Quick Demo)**

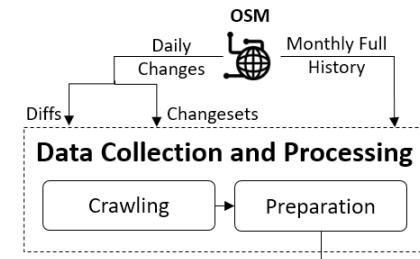
# RASED Architecture



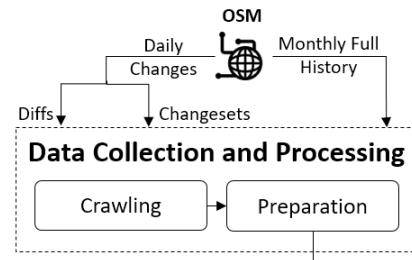
# RASED Architecture



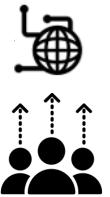
# 1- Data Collection and Processing



# 1- Data Collection and Processing



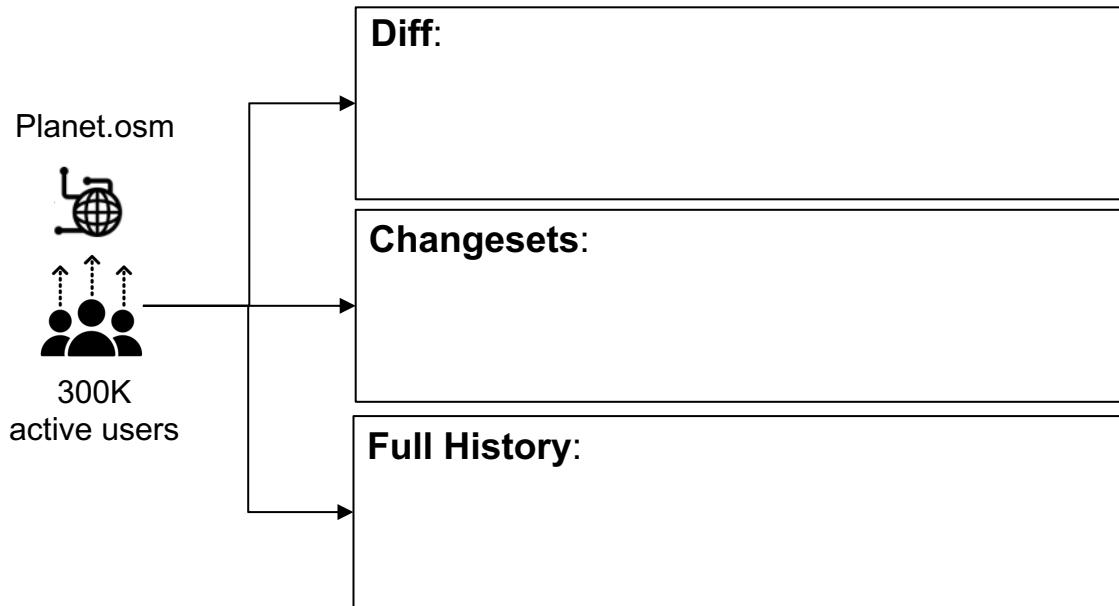
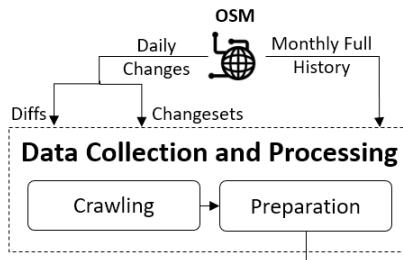
Planet.osm



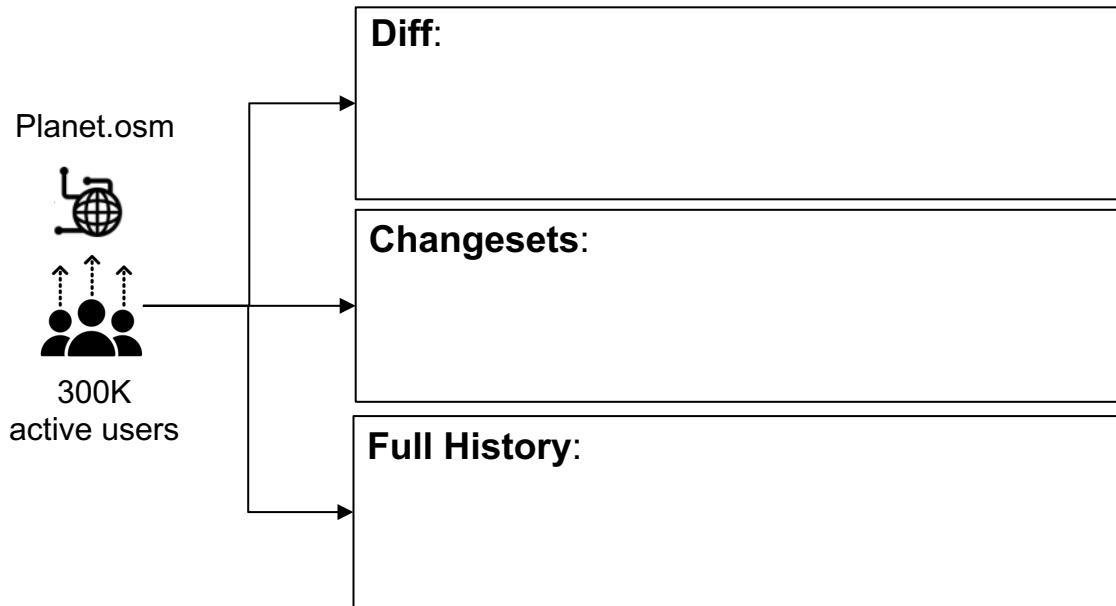
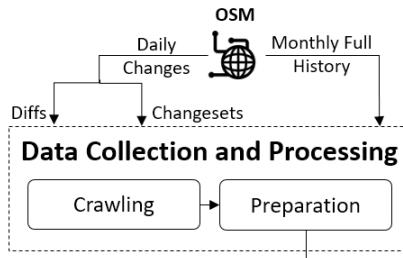
300K

active users

# 1- Data Collection and Processing



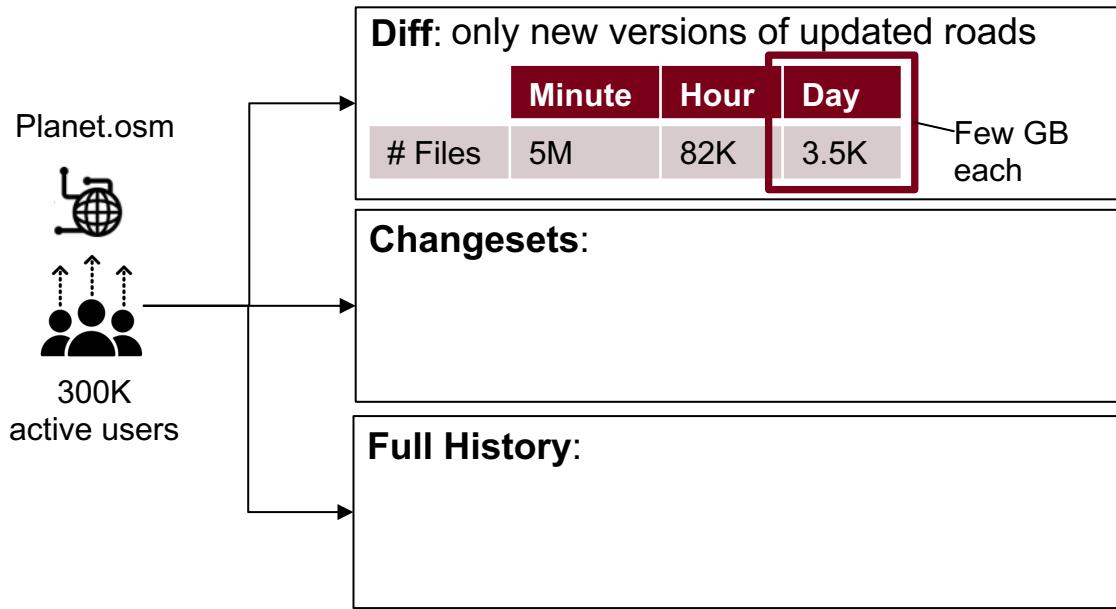
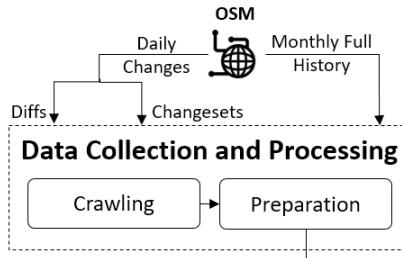
# 1- Data Collection and Processing



UpdateList U

Date	Country	RoadType	ElementType	UpdateType	Lat	Long	ChangesetID
?	?	?	?	?	?	?	?

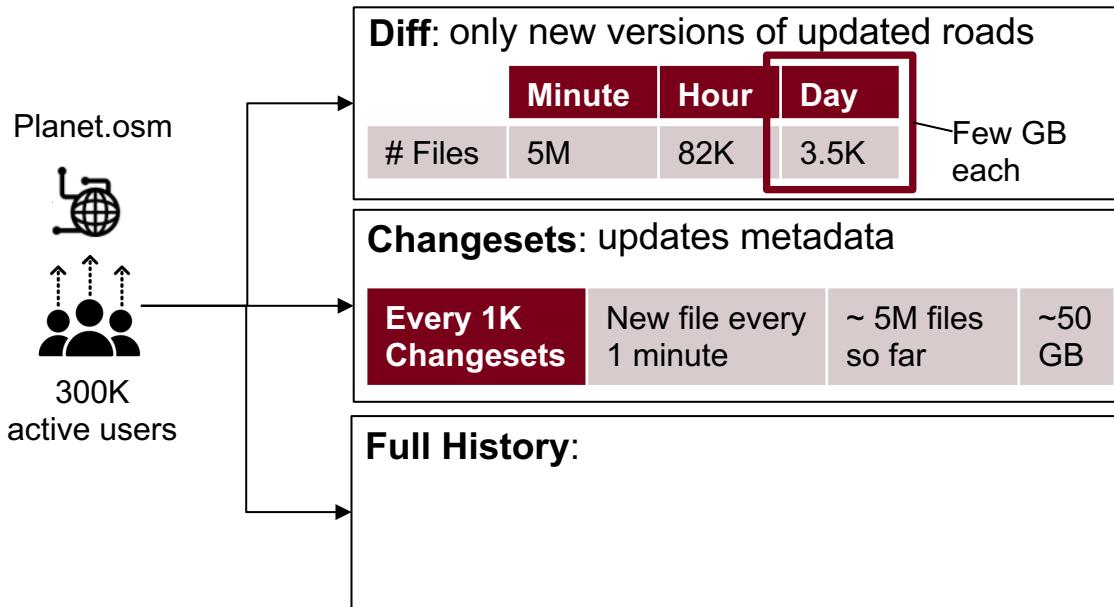
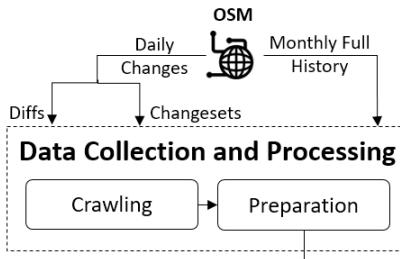
# 1- Data Collection and Processing



UpdateList U

Date	Country	RoadType	ElementType	UpdateType	Lat	Long	ChangesetID
?	?	?	?	?	?	?	?

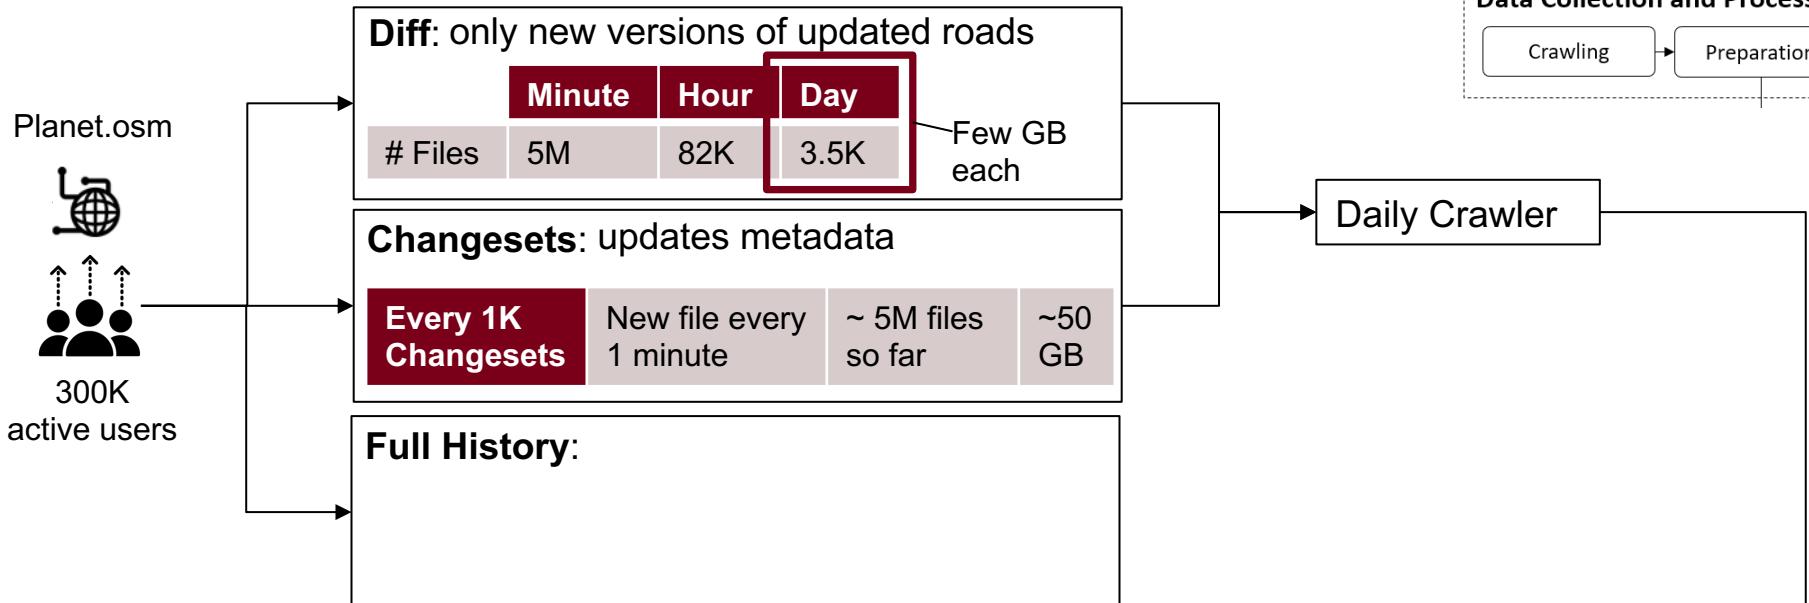
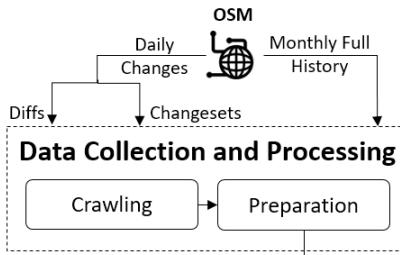
# 1- Data Collection and Processing



UpdateList U

Date	Country	RoadType	ElementType	UpdateType	Lat	Long	ChangesetID
?	?	?	?	?	?	?	?

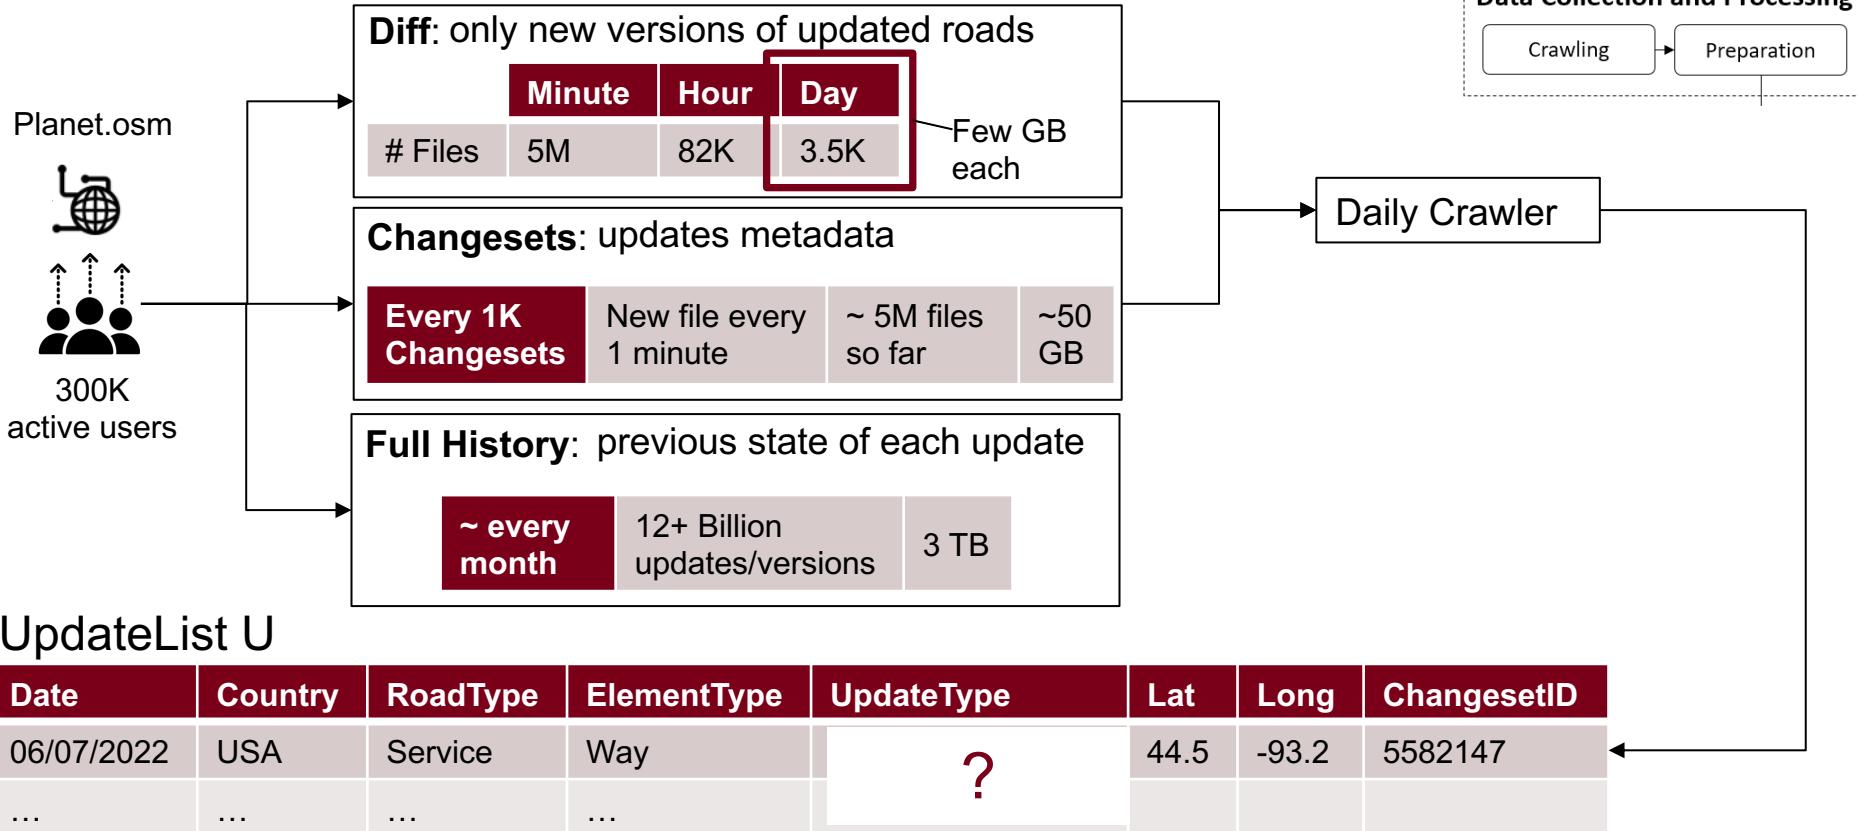
# 1- Data Collection and Processing



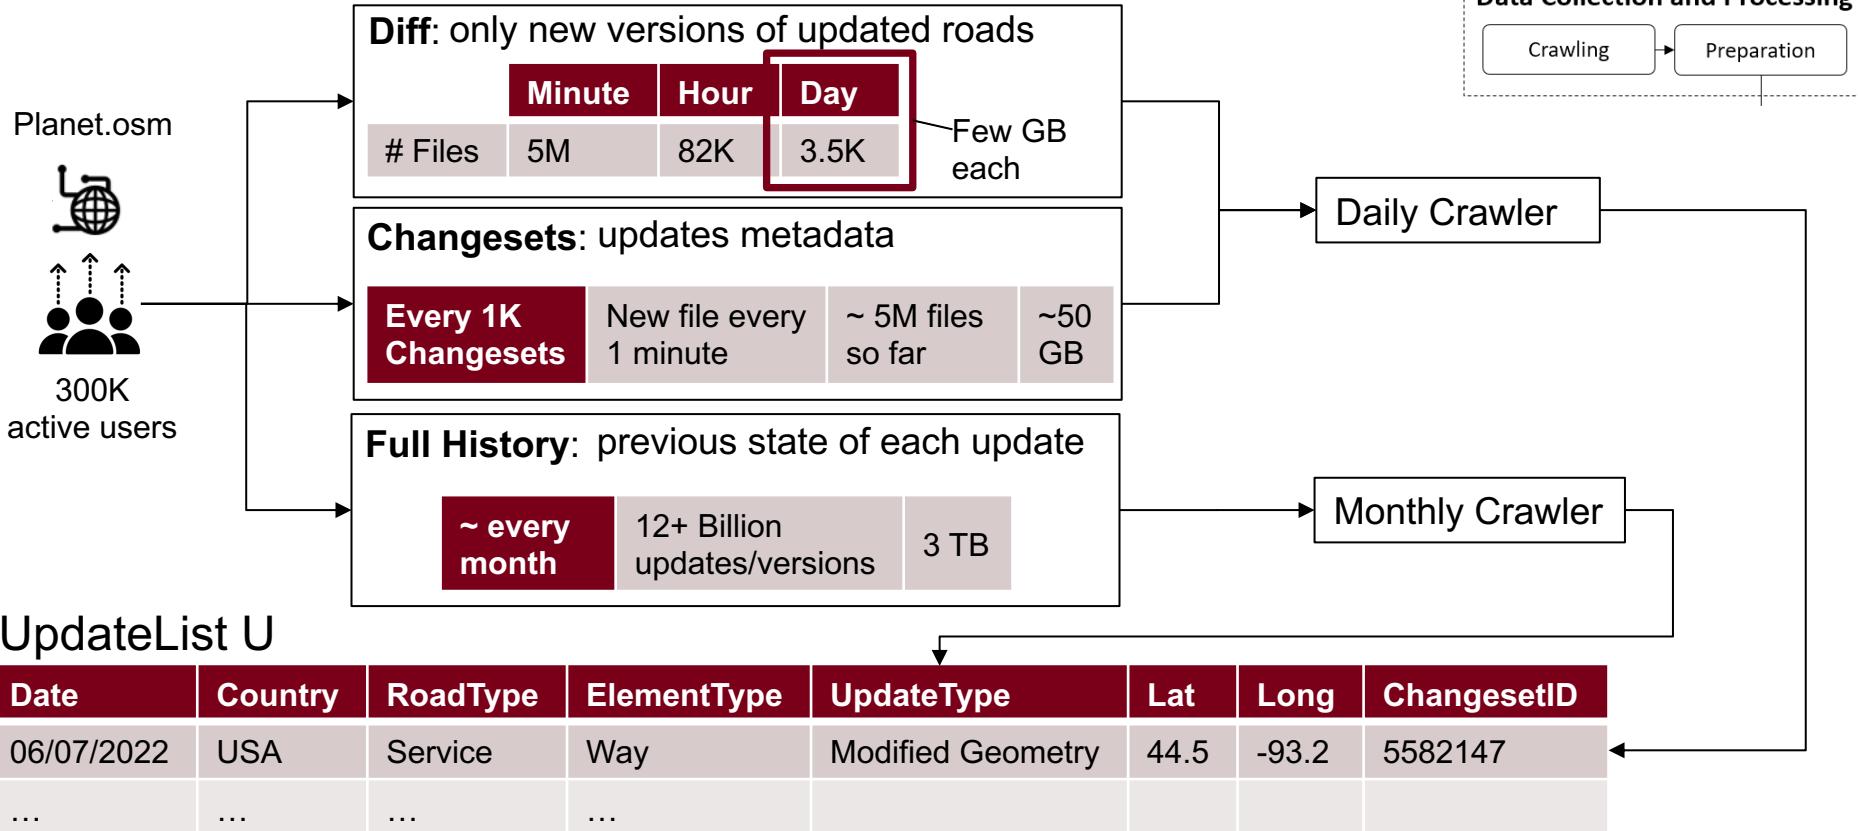
UpdateList U

Date	Country	RoadType	ElementType	UpdateType	Lat	Long	ChangesetID
06/07/2022	USA	Service	Way	?	44.5	-93.2	5582147
...	...	...	...				

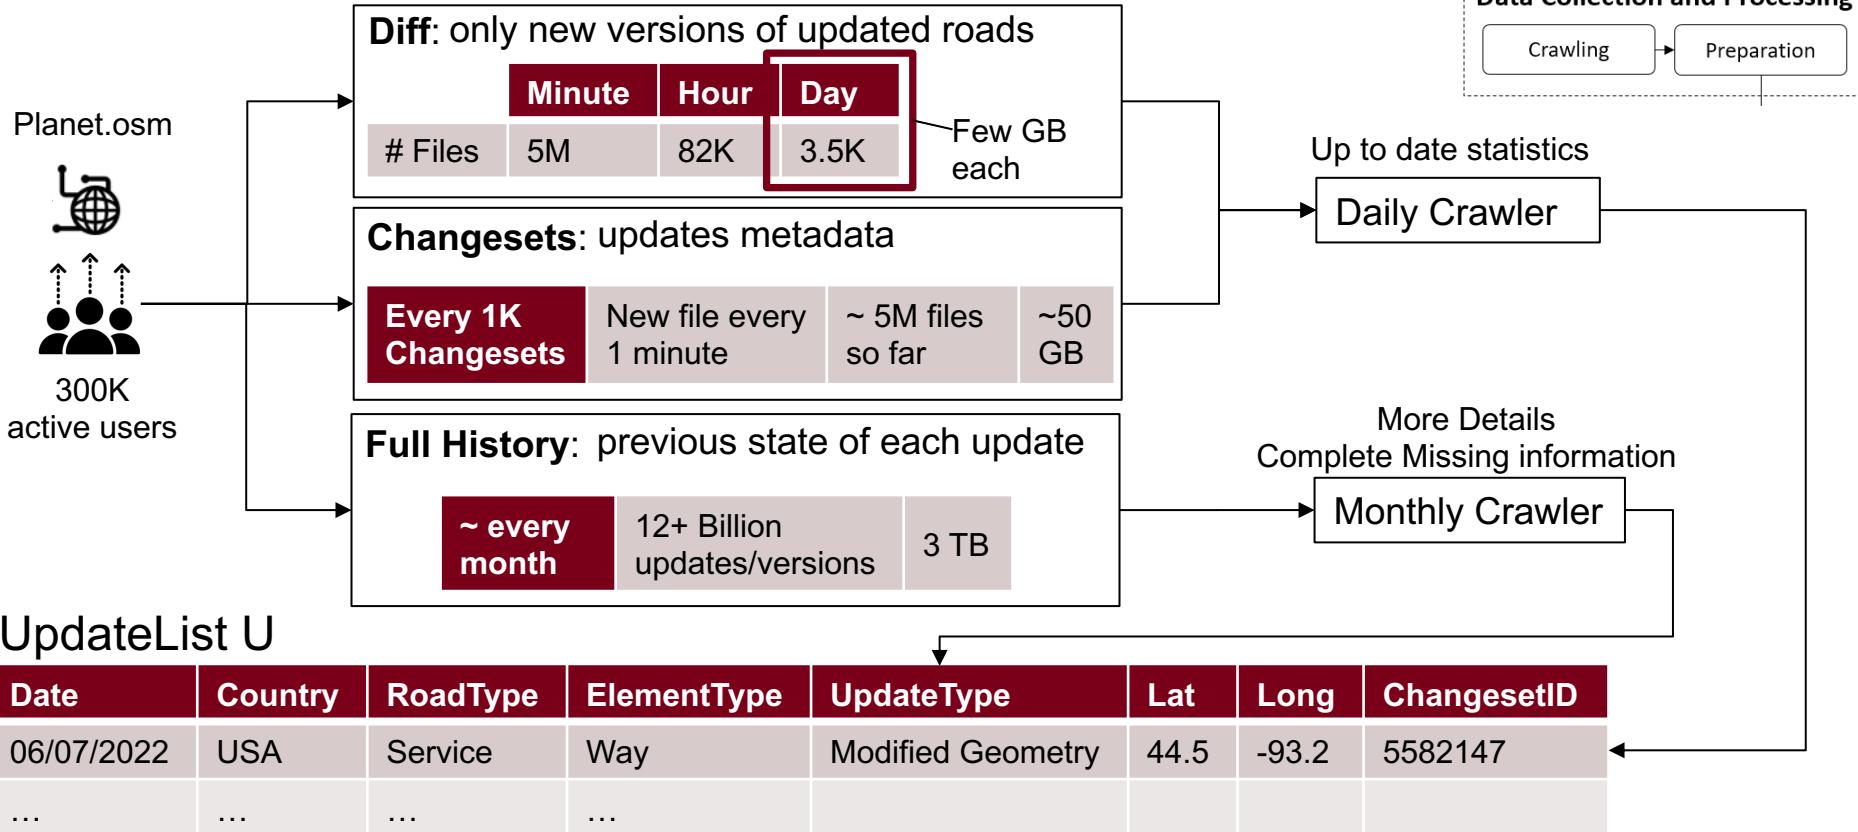
# 1- Data Collection and Processing



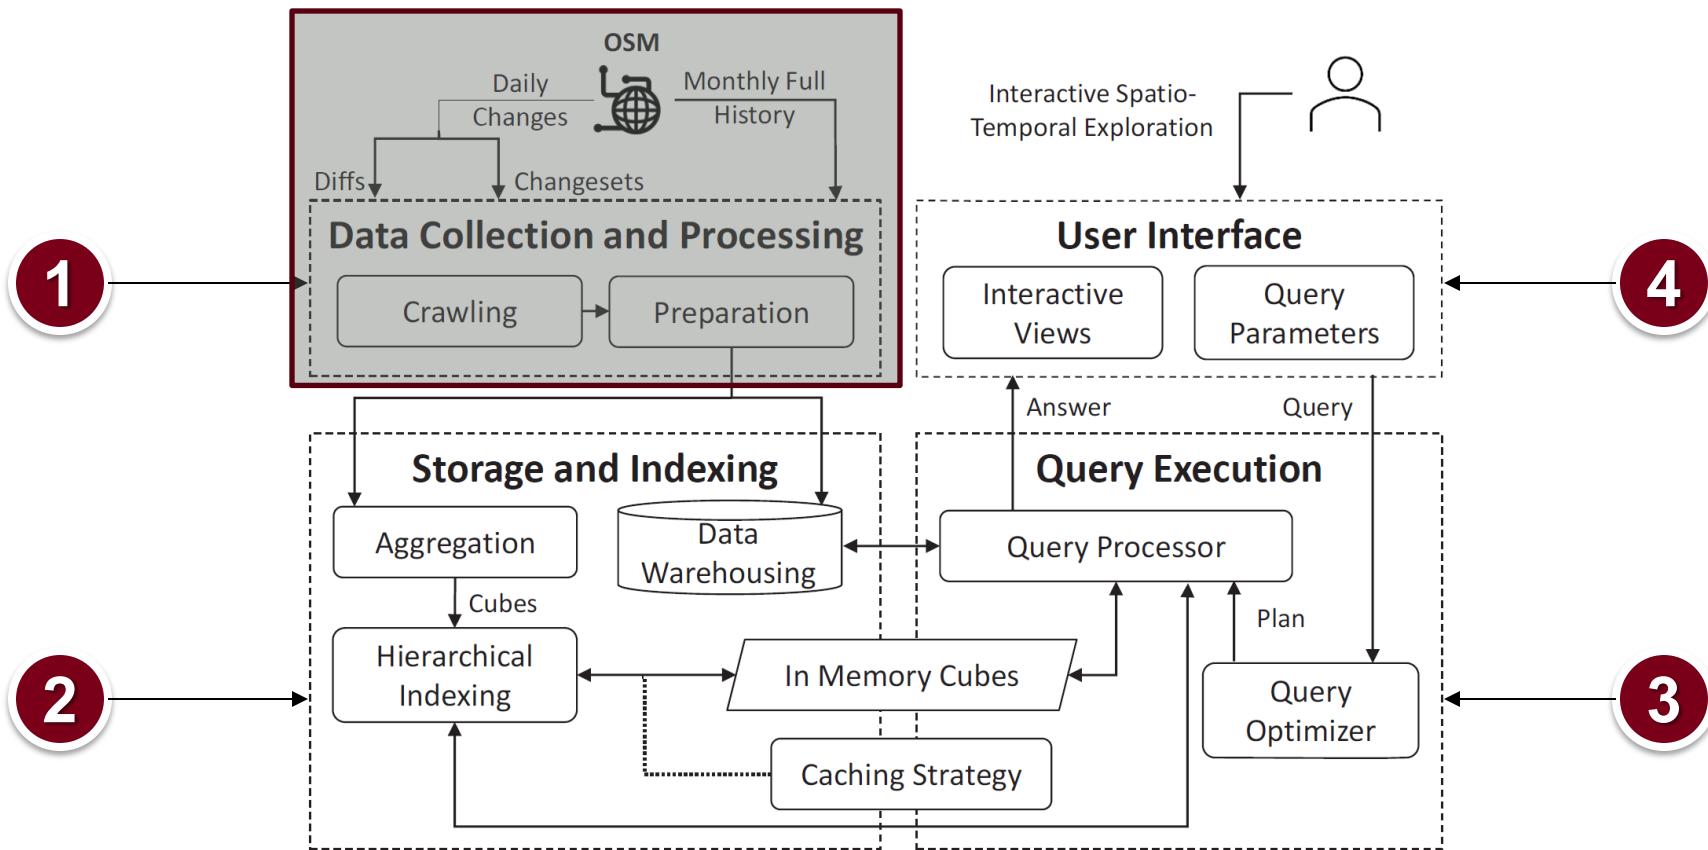
# 1- Data Collection and Processing



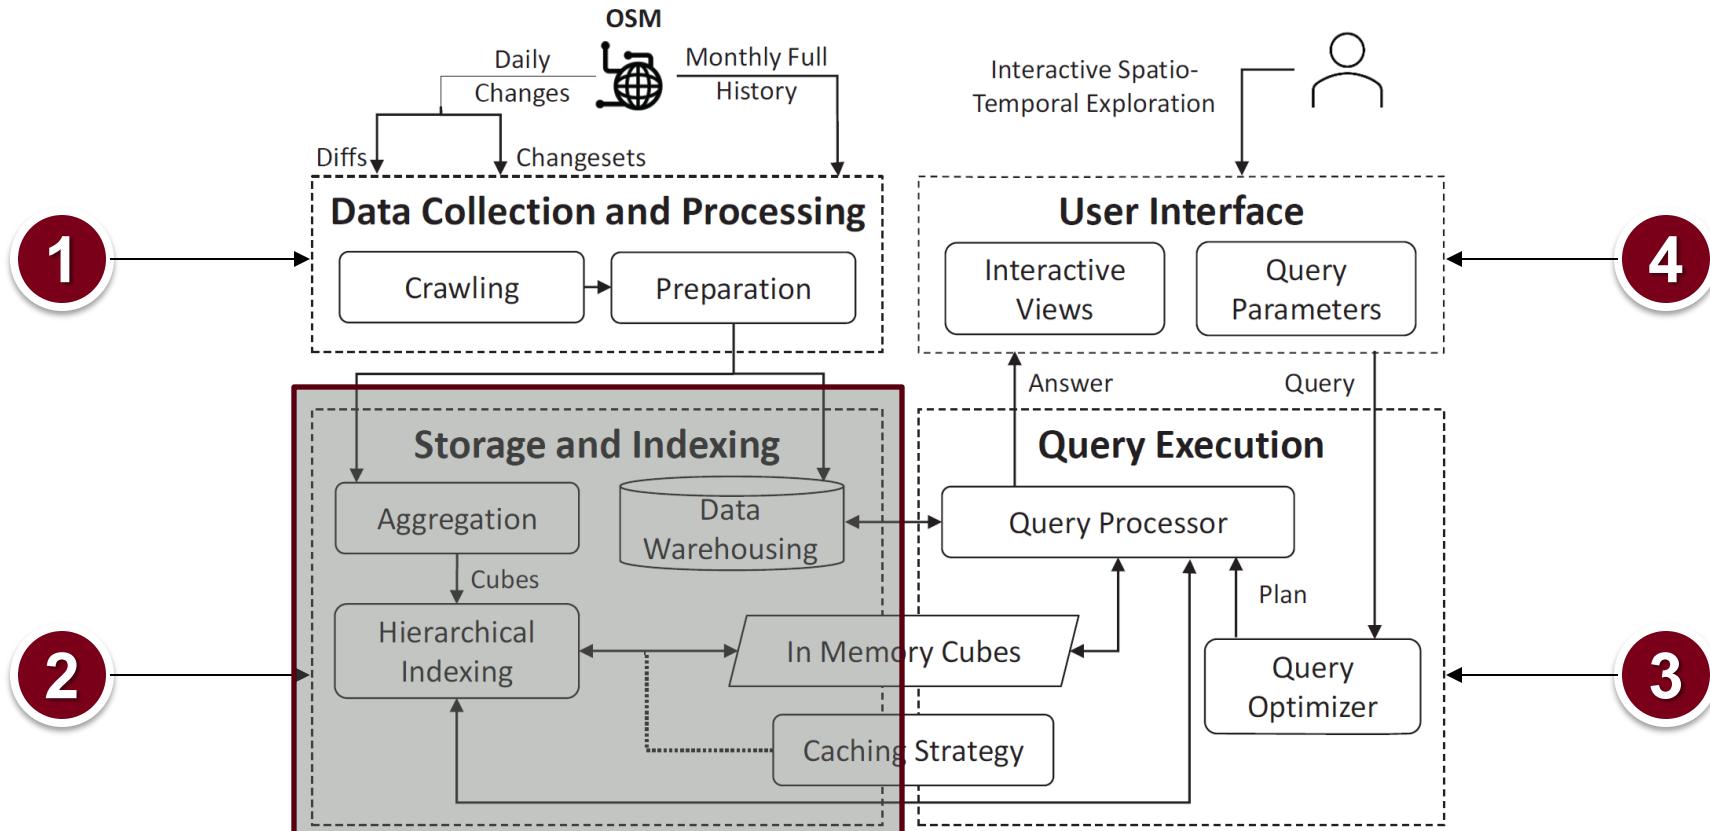
# 1- Data Collection and Processing



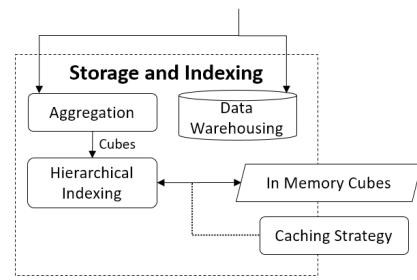
# RASED Architecture



# RASED Architecture

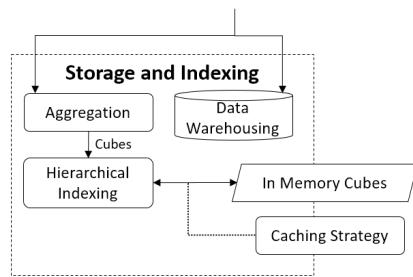


# 2- Storage and Indexing



# 2- Storage and Indexing

Daily Update List

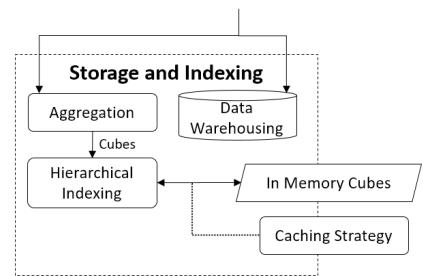


# 2- Storage and Indexing

Daily Update List

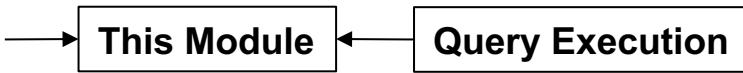


Effective Storage  
Infrastructure



# 2- Storage and Indexing

Daily Update List



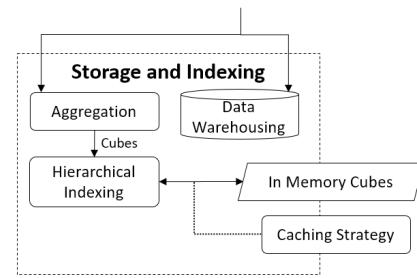
Effective Storage Infrastructure

**Very few I/O**

**Query Execution**

RASED Analytical Queries:

- Global-Scale
- Comprehensive
- Highly Interactive



# 2- Storage and Indexing



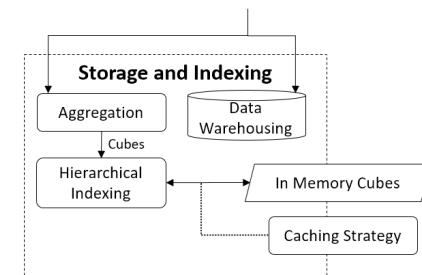
Effective Storage Infrastructure

Very few I/O

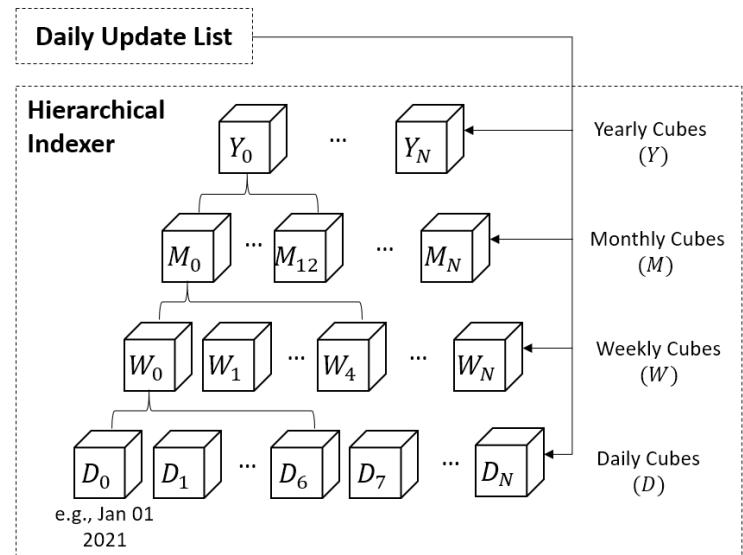
**Query Execution**

RASED Analytical Queries:

- Global-Scale
- Comprehensive
- Highly Interactive



## Hierarchical Temporal Index



# 2- Storage and Indexing



Effective Storage Infrastructure

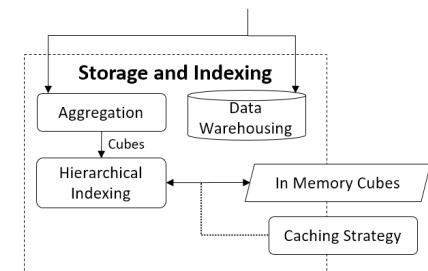
Very few I/O

RASED Analytical Queries:

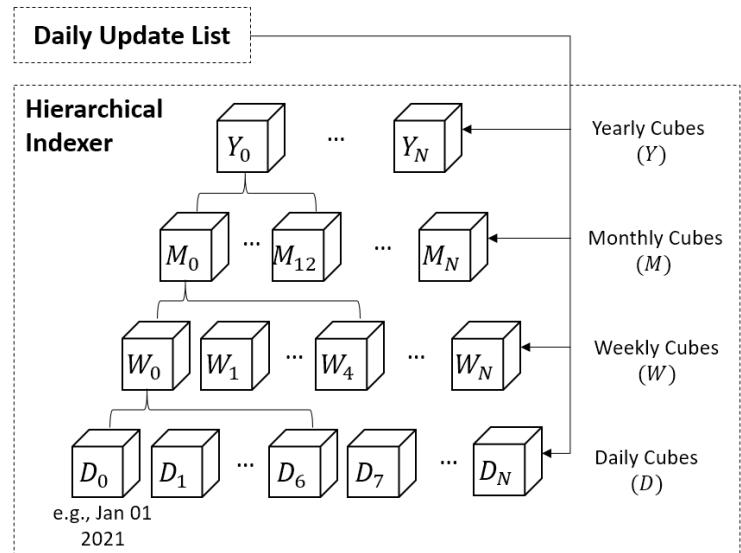
- Global-Scale
- Comprehensive
- Highly Interactive

## Index Nodes

- Precomputed statistics, not map updates itself
- Stored as 4-dimentional data cubes:
  - **Country**: 300+ values
  - **Road Type**: 150 values
  - **Update Type**: 4 values
  - **Element Type**: 3 values.



## Hierarchical Temporal Index



# 2- Storage and Indexing



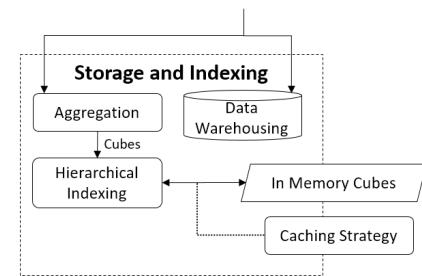
Effective Storage Infrastructure

Very few I/O

**Query Execution**

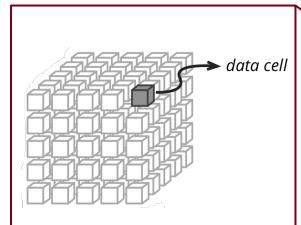
RASED Analytical Queries:

- Global-Scale
- Comprehensive
- Highly Interactive

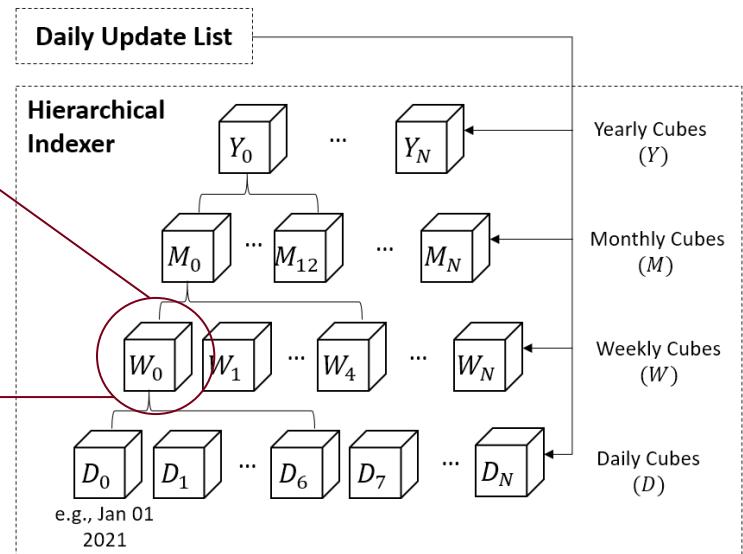


## Index Nodes

- Precomputed statistics, not map updates itself
- Stored as 4-dimensional data cubes:
  - **Country**: 300+ values
  - **Road Type**: 150 values
  - **Update Type**: 4 values
  - **Element Type**: 3 values.
- Each cube:
  - has 540,000 precomputed values (cells)
  - covers a specific period (e.g., week)
  - ~4MB, one disk page.
- Each cube cell is the aggregate count of OSM updates

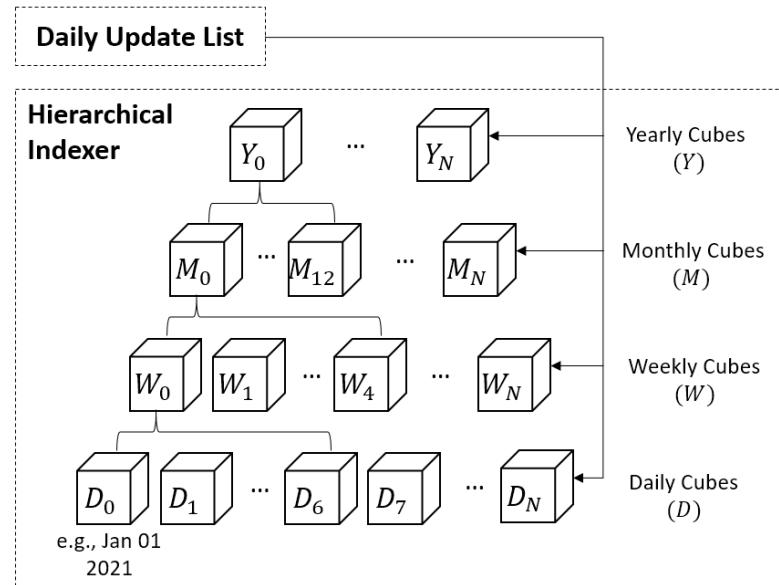


## Hierarchical Temporal Index



# RASED Hierarchical Temporal Indexing

## Hierachal Temporal Index



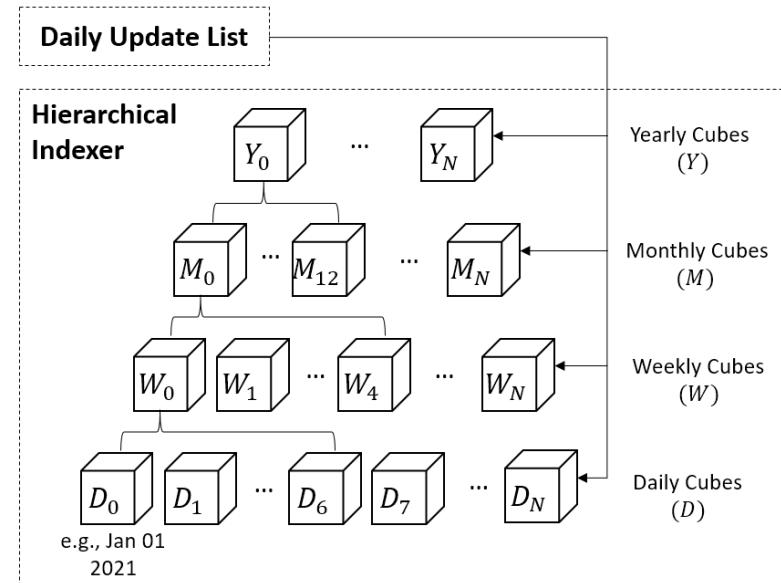
# RASED Hierarchical Temporal Indexing

## Index Hierarchy

Four levels:

- yearly
- monthly
- weekly
- daily

## Hierachal Temporal Index



# RASED Hierarchical Temporal Indexing

## Index Hierarchy

Four levels:

- yearly
- monthly
- weekly
- daily

## Index Size

OSM updates since 2004:

6,000+ daily nodes

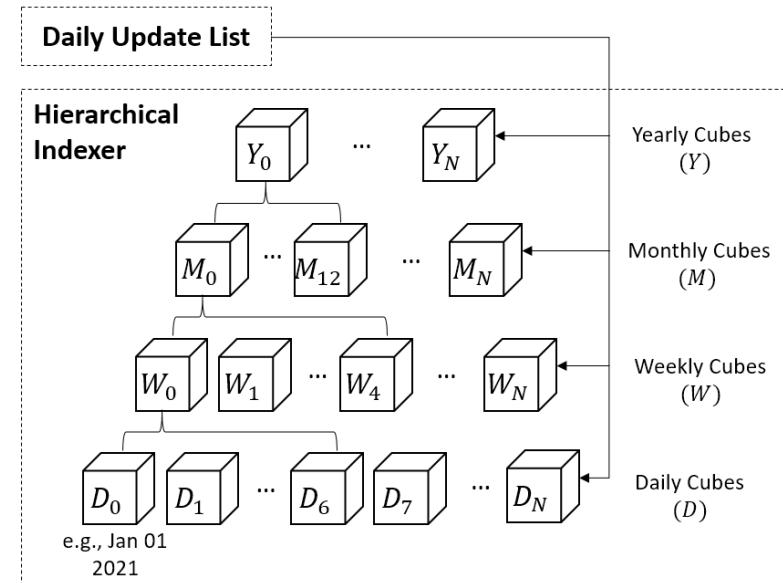
200+ monthly nodes

850+ weekly nodes

16 yearly nodes

**7,000 nodes / 4 billion aggregate values / ~28GB**

## Hierachal Temporal Index



# RASED Hierarchical Temporal Indexing

## Index Hierarchy

Four levels:

- yearly
- monthly
- weekly
- daily

OSM updates since 2004:

6,000+ daily nodes  
200+ monthly nodes

850+ weekly nodes  
16 yearly nodes

**7,000 nodes / 4 billion aggregate values / ~28GB**

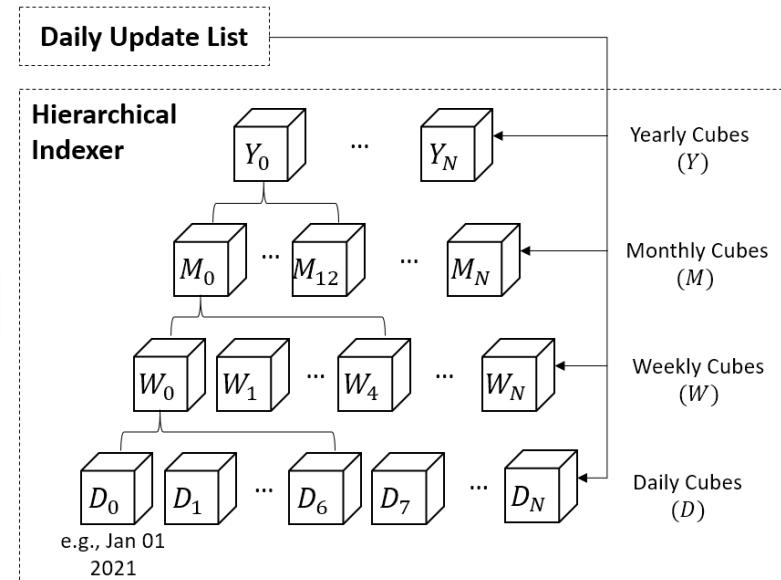
## Index Size

## Index Maintenance

### Daily updates :

- every day constructs & appends new daily cube
- Recursively build weekly/monthly/yearly cubes
- offline, ~30 minutes

## Hierachal Temporal Index



# RASED Hierarchical Temporal Indexing

## Index Hierarchy

Four levels:

- yearly
- monthly
- weekly
- daily

OSM updates since 2004:

6,000+ daily nodes  
200+ monthly nodes

850+ weekly nodes  
16 yearly nodes

**7,000 nodes / 4 billion aggregate values / ~28GB**

## Index Size

## Index Maintenance

### Daily updates :

- every day constructs & appends new daily cube
- Recursively build weekly/monthly/yearly cubes
- offline, ~30 minutes

### Monthly updates:

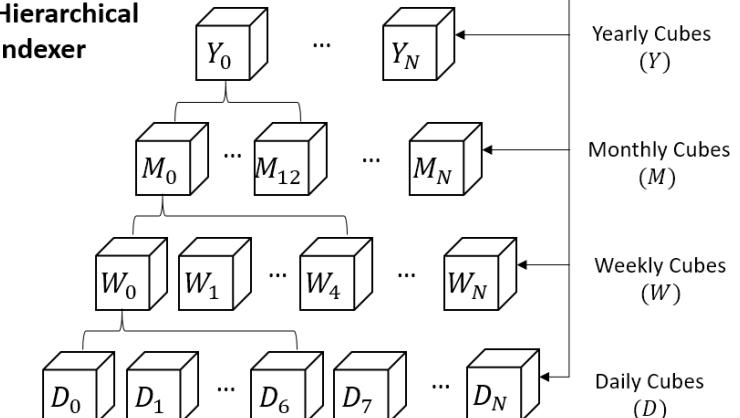
more information obtained:

- reconstruct daily & weekly cubes
- offline, ~few hours.

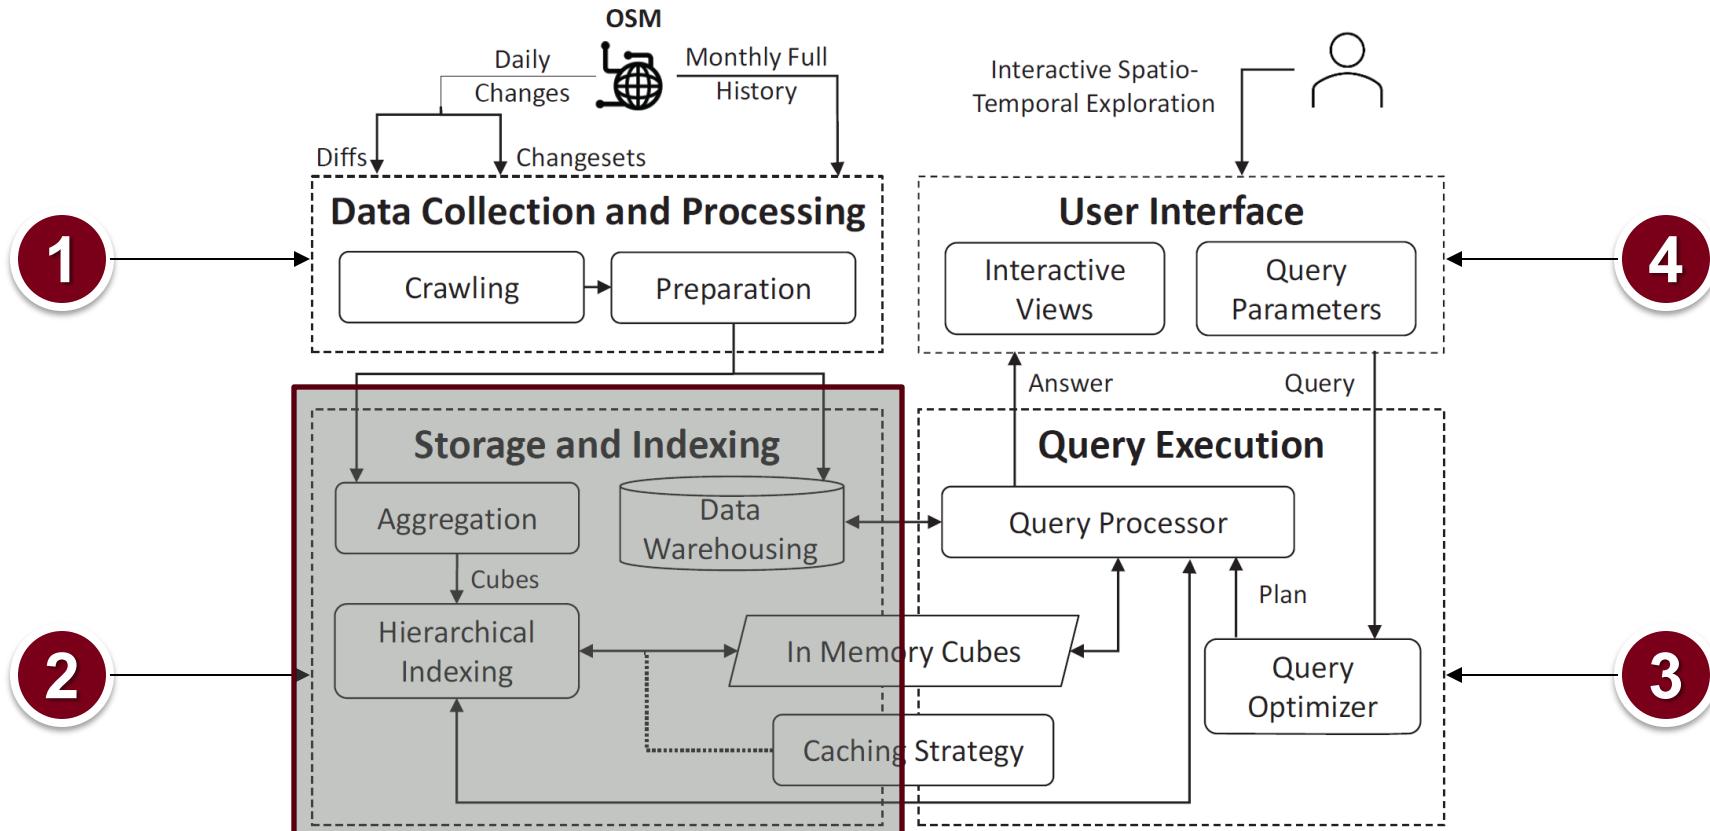
## Hierachal Temporal Index

### Daily Update List

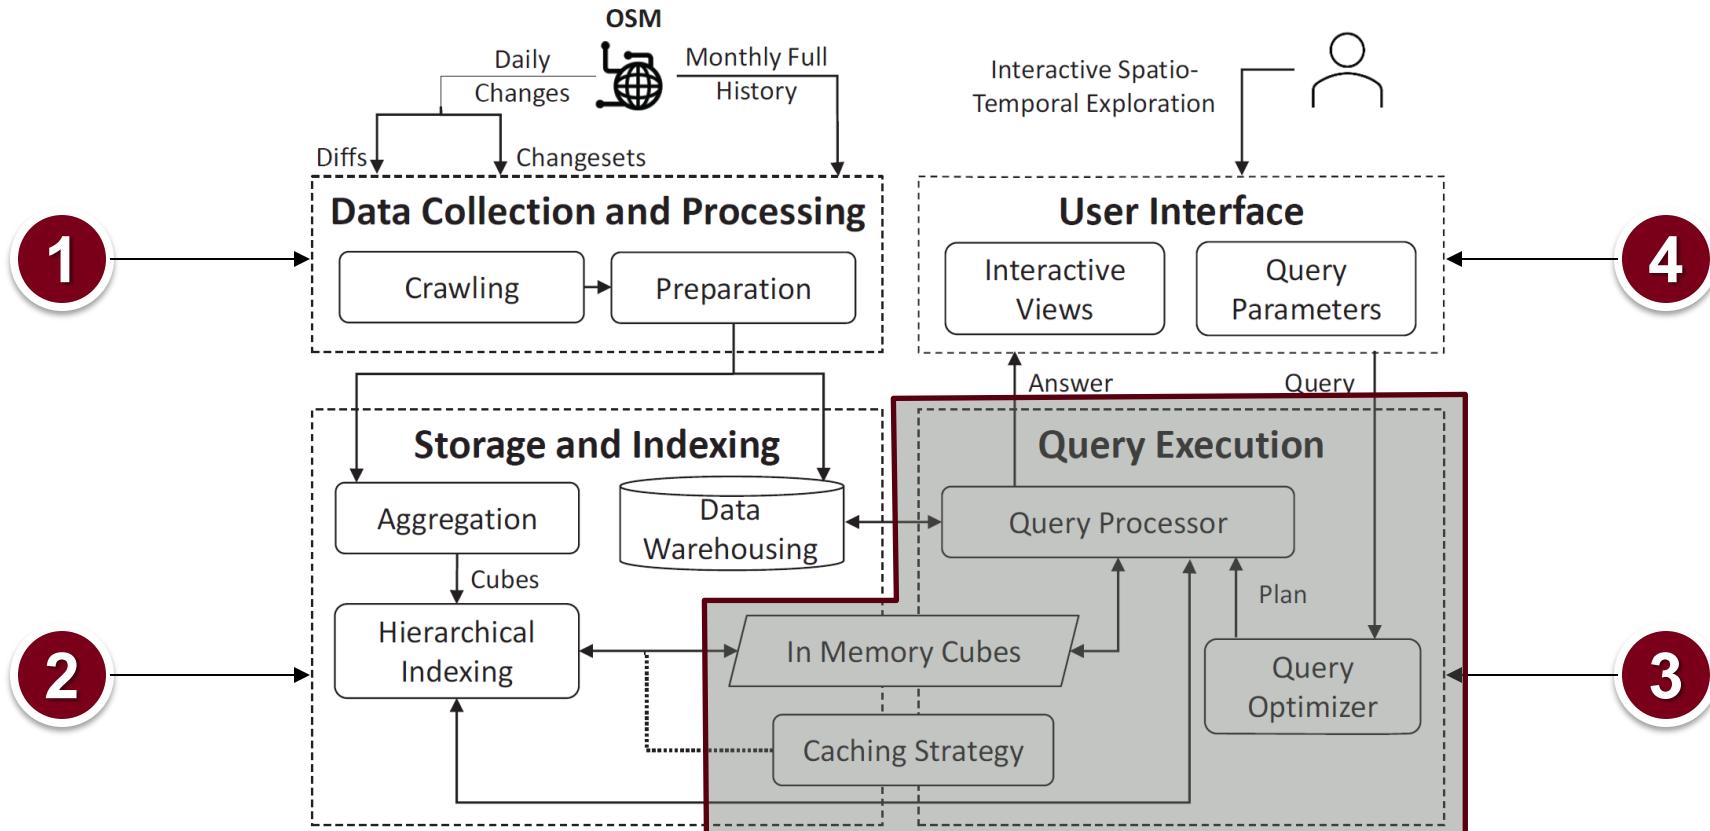
### Hierarchical Indexer



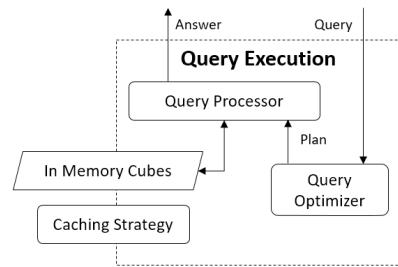
# RASED Architecture



# RASED Architecture

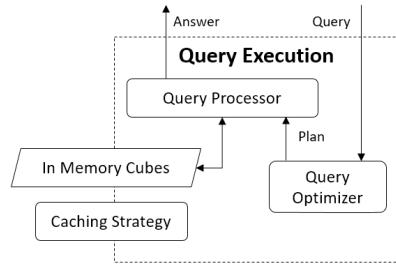


# 3- Query Execution



# 3- Query Execution

Ex. Find number of updates in each country in time window t.

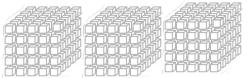


# 3- Query Execution

Ex. Find number of updates in each country in time window t.

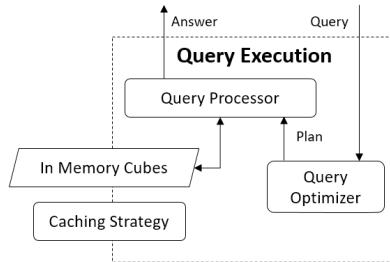
Two phases:

Retrieves cubes satisfies t



All dimensions

1

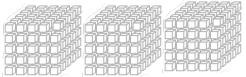


# 3- Query Execution

Ex. Find number of updates in each country in time window t.

Two phases:

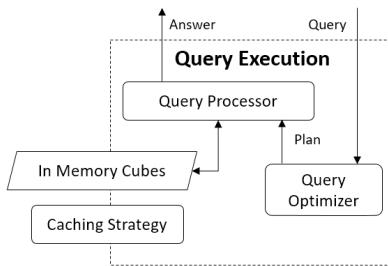
Retrieves cubes satisfies t



All dimensions

1

Disk-Based

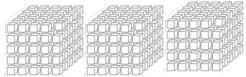


# 3- Query Execution

Ex. Find number of updates in each country in time window t.

Two phases:

Retrieves cubes satisfies t



All dimensions

1

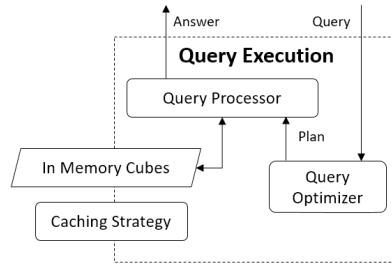
Disk-Based

Aggregates unwanted dimensions  
Keep only requested ones



Country

2

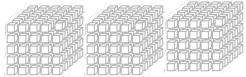


# 3- Query Execution

Ex. Find number of updates in each country in time window t.

Two phases:

Retrieves cubes satisfies t



All dimensions

1

Disk-Based

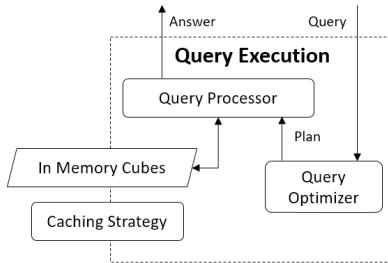
Aggregates unwanted dimensions  
Keep only requested ones



Country

2

In-Memory

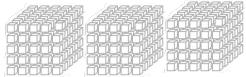


# 3- Query Execution

Ex. Find number of updates in each country in time window t.

Two phases:

Retrieves cubes satisfies t



All dimensions

1

Disk-Based

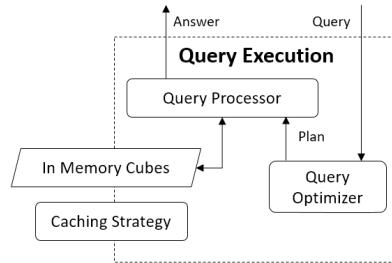
Aggregates unwanted dimensions  
Keep only requested ones



Country

2

In-Memory

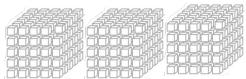


# 3- Query Execution

Ex. Find number of updates in each country in time window t.

Two phases:

Retrieves cubes satisfies t



All dimensions

1

Disk-Based

Aggregates unwanted dimensions  
Keep only requested ones



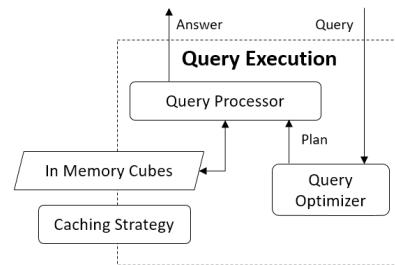
Country

2

In-Memory

Two Disk Access optimization techniques

Caching & Level Optimization

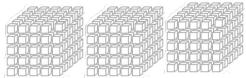


# 3- Query Execution

Ex. Find number of updates in each country in time window t.

Two phases:

Retrieves cubes satisfies t



All dimensions

1

Disk-Based

Aggregates unwanted dimensions  
Keep only requested ones



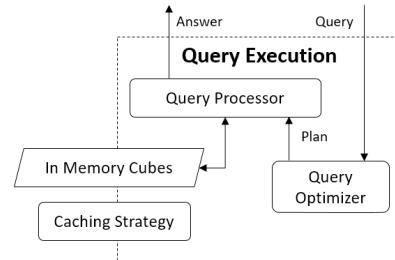
Country

2

In-Memory

Two Disk Access optimization techniques

## Caching & Level Optimization



### A. Caching

Goal:

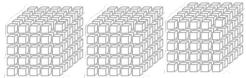
- Preload recent data cubes into memory
- Reduce cubes retrieved from disk
- Reduce query response time

# 3- Query Execution

Ex. Find number of updates in each country in time window t.

Two phases:

Retrieves cubes satisfies t



All dimensions

1

Disk-Based

Aggregates unwanted dimensions  
Keep only requested ones



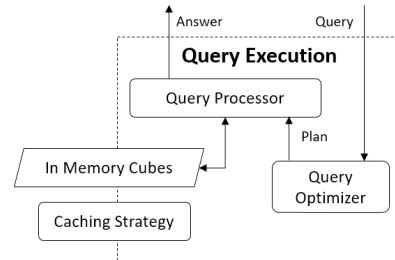
Country

2

In-Memory

Two Disk Access optimization techniques

## Caching & Level Optimization



### A. Caching

#### Goal:

- Preload recent data cubes into memory
- Reduce cubes retrieved from disk
- Reduce query response time

#### Rational:

More likely to receive queries for recent updates

#### Challenge:

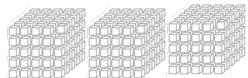
which level to preloaded data cubes from?

# 3- Query Execution

Ex. Find number of updates in each country in time window t.

Two phases:

Retrieves cubes satisfies t



All dimensions

1

Disk-Based

Aggregates unwanted dimensions  
Keep only requested ones



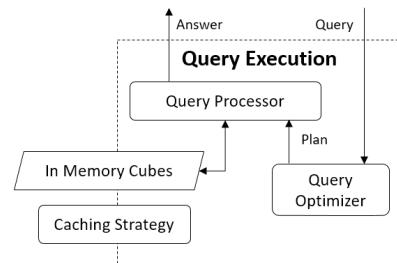
Country

2

In-Memory

Two Disk Access optimization techniques

## Caching & Level Optimization



### A. Caching

#### Goal:

- Preload recent data cubes into memory
- Reduce cubes retrieved from disk
- Reduce query response time

#### Rational:

More likely to receive queries for recent updates

#### Challenge:

which level to preloaded data cubes from?

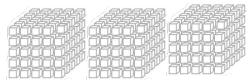
$$\{D_{|D|-i}\}_{i=0}^{\alpha N} \cup \{W_{|W|-i}\}_{i=0}^{\beta N} \cup \{M_{|M|-i}\}_{i=0}^{\gamma N} \cup \{Y_{|Y|-i}\}_{i=0}^{\theta N}$$

# 3- Query Execution

Ex. Find number of updates in each country in time window t.

Two phases:

Retrieves cubes satisfies t



All dimensions

1

Disk-Based

Aggregates unwanted dimensions  
Keep only requested ones



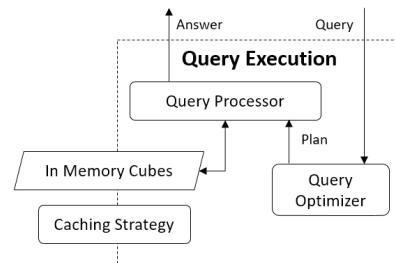
Country

2

In-Memory

Two Disk Access optimization techniques

## Caching & Level Optimization



### A. Caching

**Goal:**

- Preload recent data cubes into memory
- Reduce cubes retrieved from disk
- Reduce query response time

**Rational:**

More likely to receive queries for recent updates

**Challenge:**

which level to preloaded data cubes from?

$$\{D_{|D|-i}\}_{i=0}^{\alpha N} \cup \{W_{|W|-i}\}_{i=0}^{\beta N} \cup \{M_{|M|-i}\}_{i=0}^{\gamma N} \cup \{Y_{|Y|-i}\}_{i=0}^{\theta N}$$

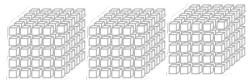
- N available memory
- $\alpha, \beta, \gamma$ , and  $\theta$ : ratio allocated to each level
- trade off: *aggregation granularity* vs *time coverage*

# 3- Query Execution

Ex. Find number of updates in each country in time window t.

Two phases:

Retrieves cubes satisfies t



All dimensions

1

Disk-Based

Aggregates unwanted dimensions  
Keep only requested ones



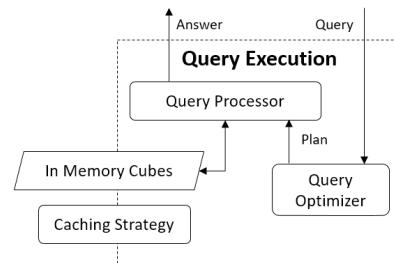
Country

2

In-Memory

Two Disk Access optimization techniques

## Caching & Level Optimization



### A. Caching

Goal:

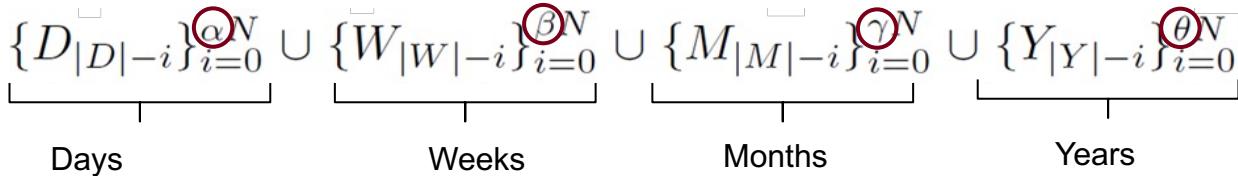
- Preload recent data cubes into memory
- Reduce cubes retrieved from disk
- Reduce query response time

Rational:

More likely to receive queries for recent updates

Challenge:

which level to preloaded data cubes from?



• N available memory

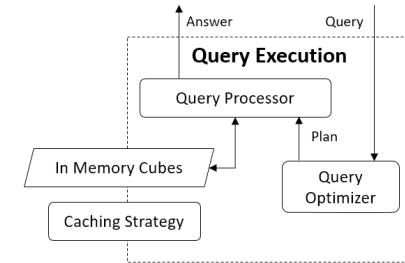
•  $\alpha, \beta, \gamma$ , and  $\theta$ : ratio allocated to each level

• trade off: *aggregation granularity vs time coverage*

# 3- Query Execution

Two Disk Access optimization techniques

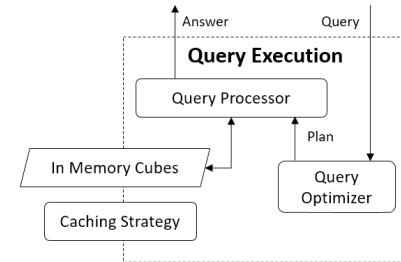
## Caching & Level Optimization



# 3- Query Execution

Two Disk Access optimization techniques

## Caching & Level Optimization



### B. Level Optimization

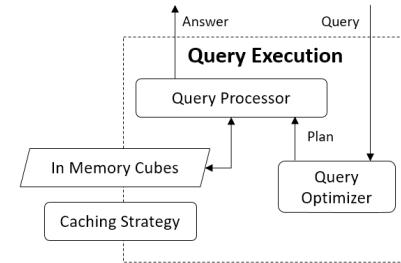
# 3- Query Execution

Two Disk Access optimization techniques

Example:

Find number of updates in during  
the period Jan 1 to Feb 15.  
(46 days)

## Caching & Level Optimization



### B. Level Optimization

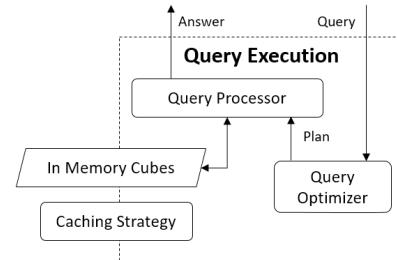
# 3- Query Execution

Two Disk Access optimization techniques

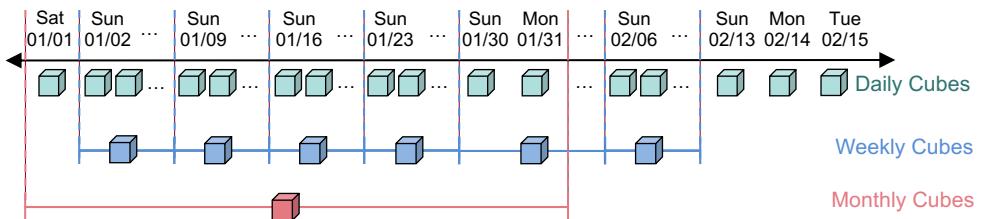
Example:

Find number of updates in during  
the period Jan 1 to Feb 15.  
(46 days)

## Caching & Level Optimization



### B. Level Optimization



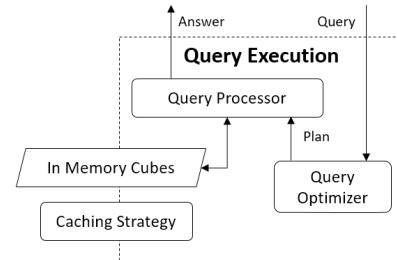
# 3- Query Execution

Two Disk Access optimization techniques

Example:

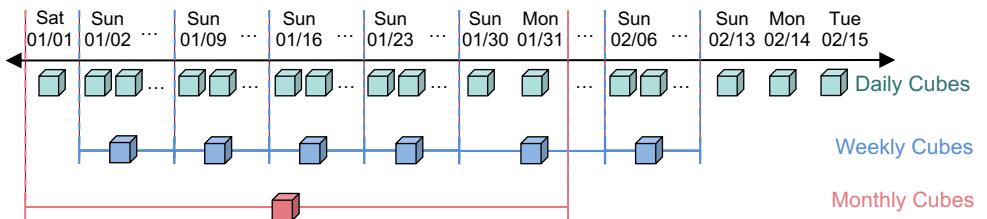
Find number of updates in during  
the period Jan 1 to Feb 15.  
(46 days)

## Caching & Level Optimization



### B. Level Optimization

Three Possible Paths:



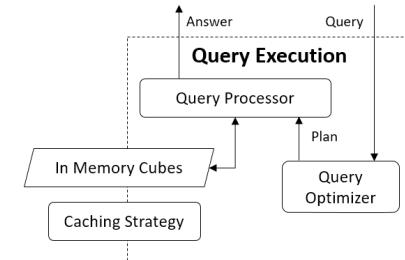
# 3- Query Execution

Two Disk Access optimization techniques

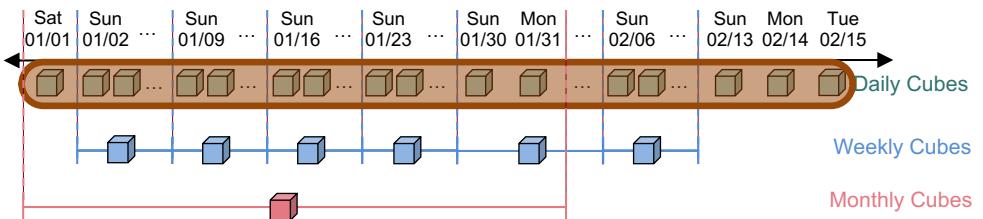
Example:

Find number of updates in during  
the period Jan 1 to Feb 15.  
(46 days)

## Caching & Level Optimization



### B. Level Optimization



Three Possible Paths:

A- All daily

cost = 46 cubes

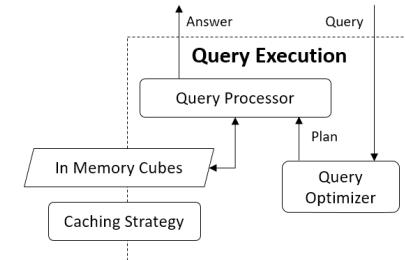
# 3- Query Execution

Two Disk Access optimization techniques

Example:

Find number of updates in during  
the period Jan 1 to Feb 15.  
(46 days)

## Caching & Level Optimization

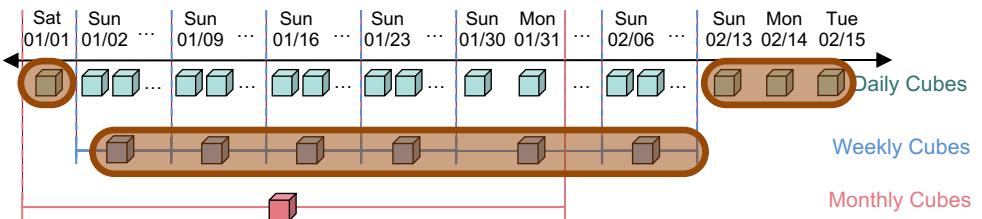


### B. Level Optimization

Three Possible Paths:

- A- All daily
- B- 6 weekly + 4 daily

cost = 46 cubes  
cost = 10 cubes



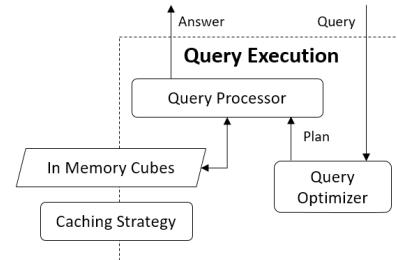
# 3- Query Execution

Two Disk Access optimization techniques

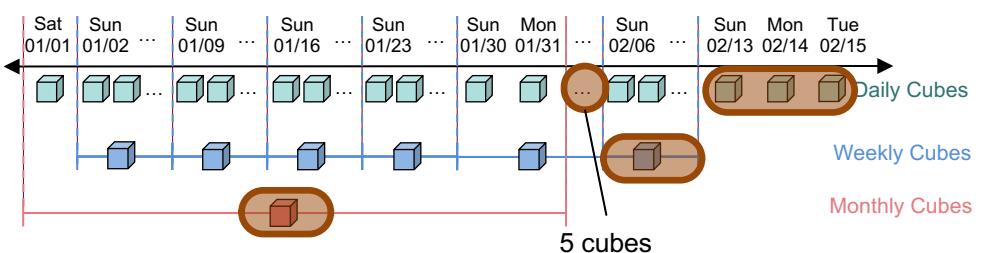
Example:

Find number of updates in during  
the period Jan 1 to Feb 15.  
(46 days)

## Caching & Level Optimization



### B. Level Optimization



Three Possible Paths:

- A- All daily
- B- 6 weekly + 4 daily
- C- 1 monthly + 1 weekly + 8 daily

cost = 46 cubes

cost = 10 cubes

cost = 10 cubes

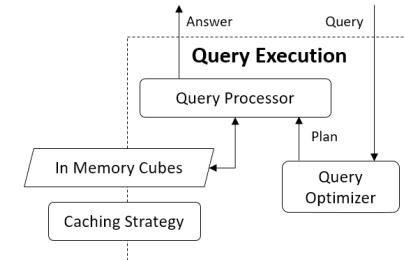
# 3- Query Execution

Two Disk Access optimization techniques

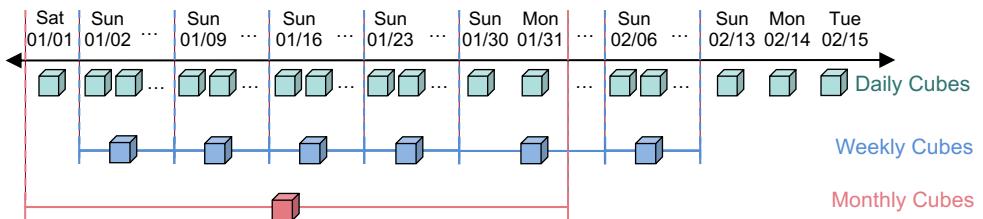
Example:

Find number of updates in during  
the period Jan 1 to Feb 15.  
(46 days)

## Caching & Level Optimization



### B. Level Optimization



Three Possible Paths:

A- All daily

cost = 46 cubes

B- 6 weekly + 4 daily

cost = 10 cubes

C- 1 monthly + 1 weekly + 8 daily

cost = 10 cubes

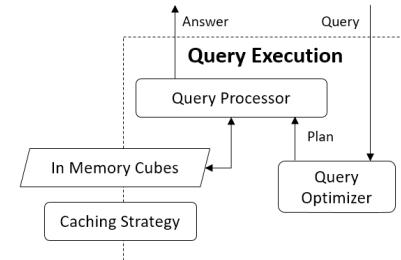
# 3- Query Execution

Two Disk Access optimization techniques

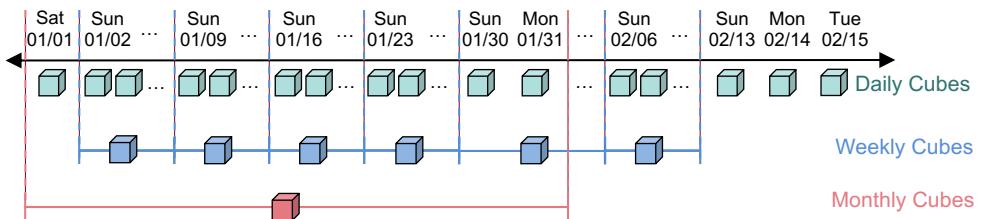
Example:

Find number of updates in during  
the period Jan 1 to Feb 15.  
(46 days)

## Caching & Level Optimization



### B. Level Optimization



Cache Awareness

Three Possible Paths:

A- All daily

cost = 46 cubes

B- 6 weekly + 4 daily

cost = 10 cubes

C- 1 monthly + 1 weekly + 8 daily

cost = 10 cubes

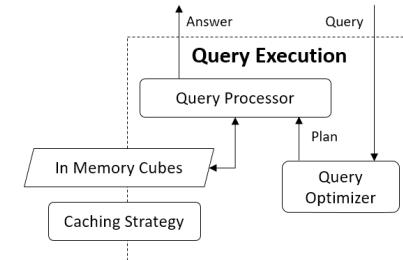
# 3- Query Execution

Two Disk Access optimization techniques

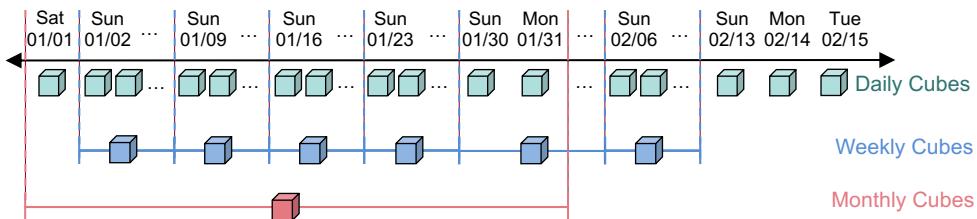
Example:

Find number of updates in during  
the period Jan 1 to Feb 15.  
(46 days)

## Caching & Level Optimization



### B. Level Optimization



Three Possible Paths:

A- All daily

cost = 46 cubes

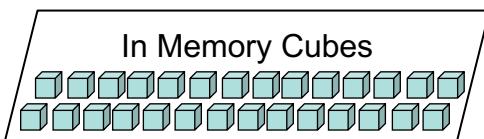
B- 6 weekly + 4 daily

cost = 10 cubes

C- 1 monthly + 1 weekly + 8 daily

cost = 10 cubes

Cache Awareness



E.g.  $\alpha = 1$

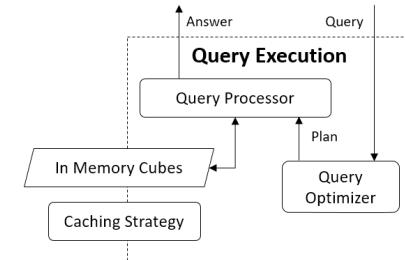
# 3- Query Execution

Two Disk Access optimization techniques

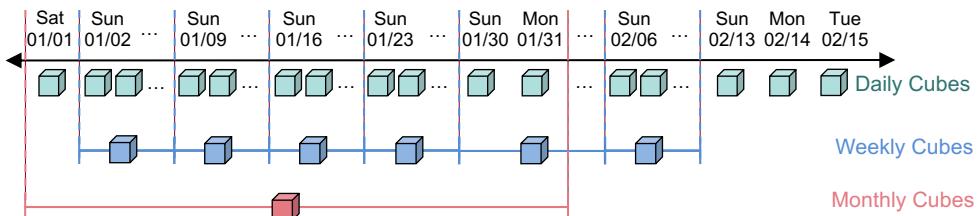
Example:

Find number of updates in during  
the period Jan 1 to Feb 15.  
(46 days)

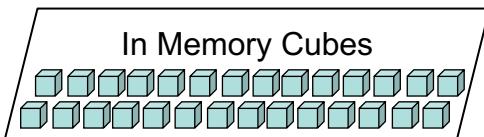
## Caching & Level Optimization



### B. Level Optimization



Cache Awareness



E.g.  $\alpha = 1$

Three Possible Paths:

A- All daily

cost = 46 cubes

B- 6 weekly + 4 daily

cost = 10 cubes

C- 1 monthly + 1 weekly + 8 daily

cost = 10 cubes

A- All daily

cost = 0 disk I/O

B- 6 weekly + 4 daily

cost = 6 disk I/O

C- 1 monthly + 1 weekly + 8 daily

cost = 2 disk I/O

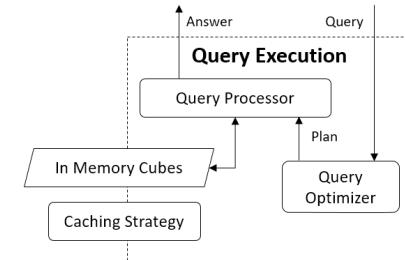
# 3- Query Execution

Two Disk Access optimization techniques

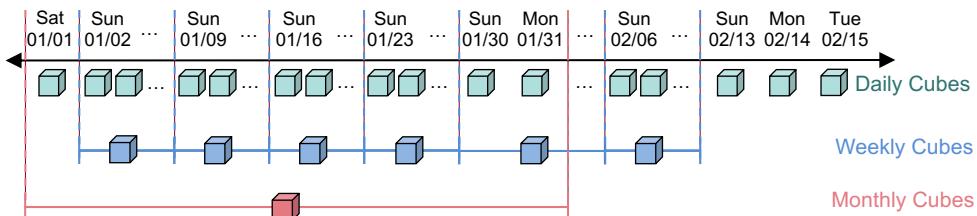
Example:

Find number of updates in during  
the period Jan 1 to Feb 15.  
(46 days)

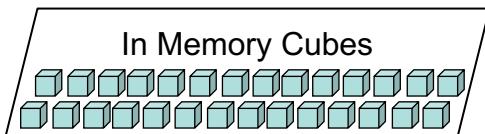
## Caching & Level Optimization



### B. Level Optimization



Cache Awareness



E.g.  $\alpha = 1$

Three Possible Paths:

- A- All daily
- B- 6 weekly + 4 daily
- C- 1 monthly + 1 weekly + 8 daily

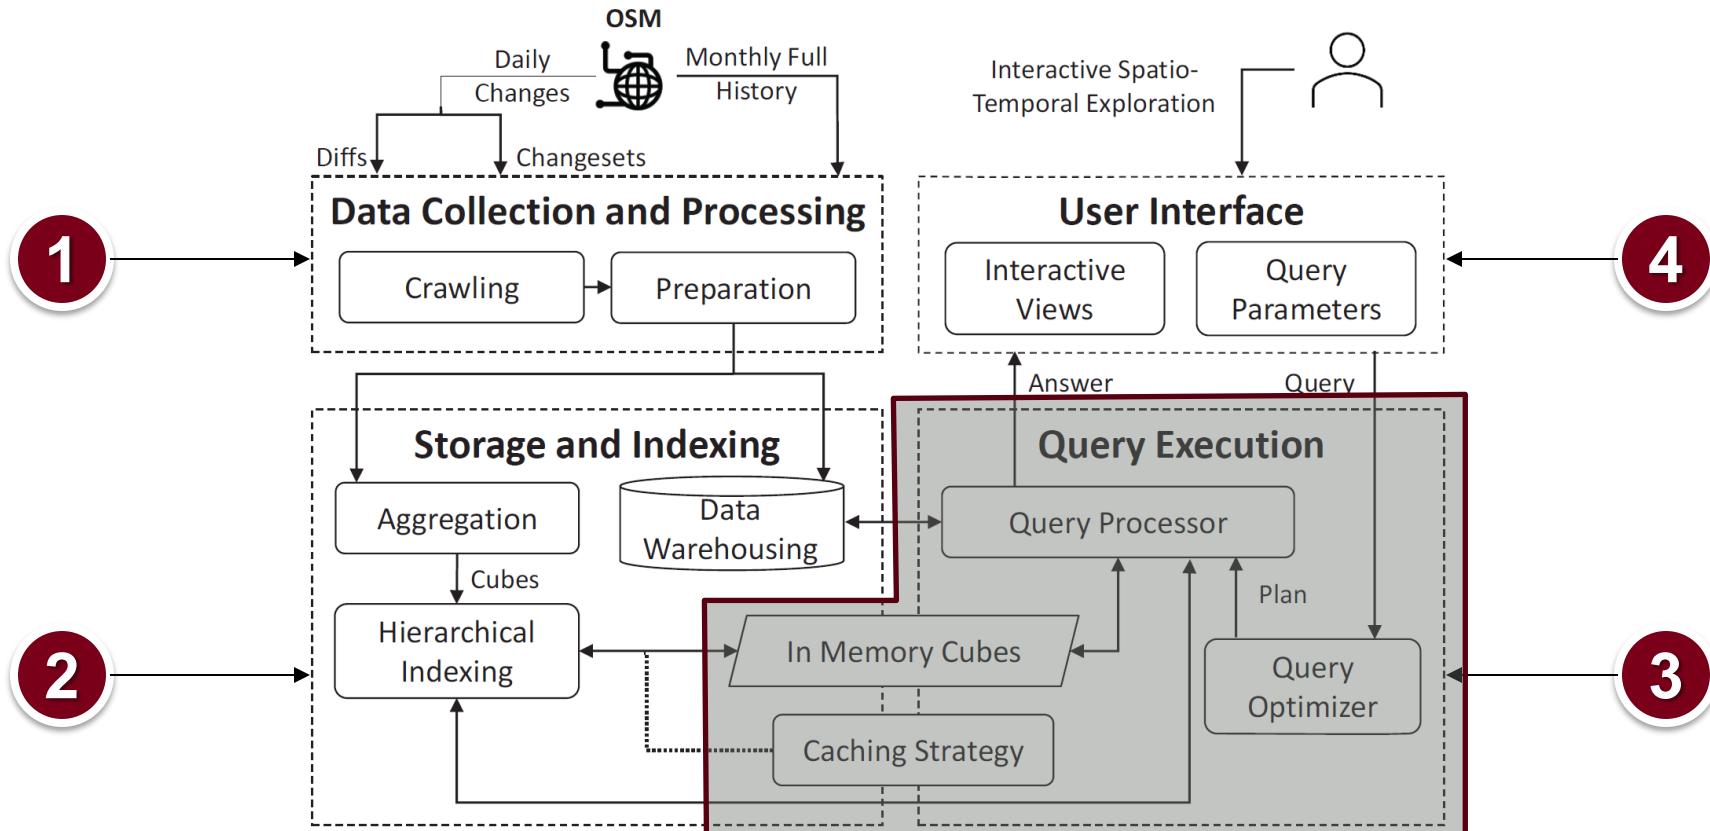
cost = 46 cubes

cost = 10 cubes

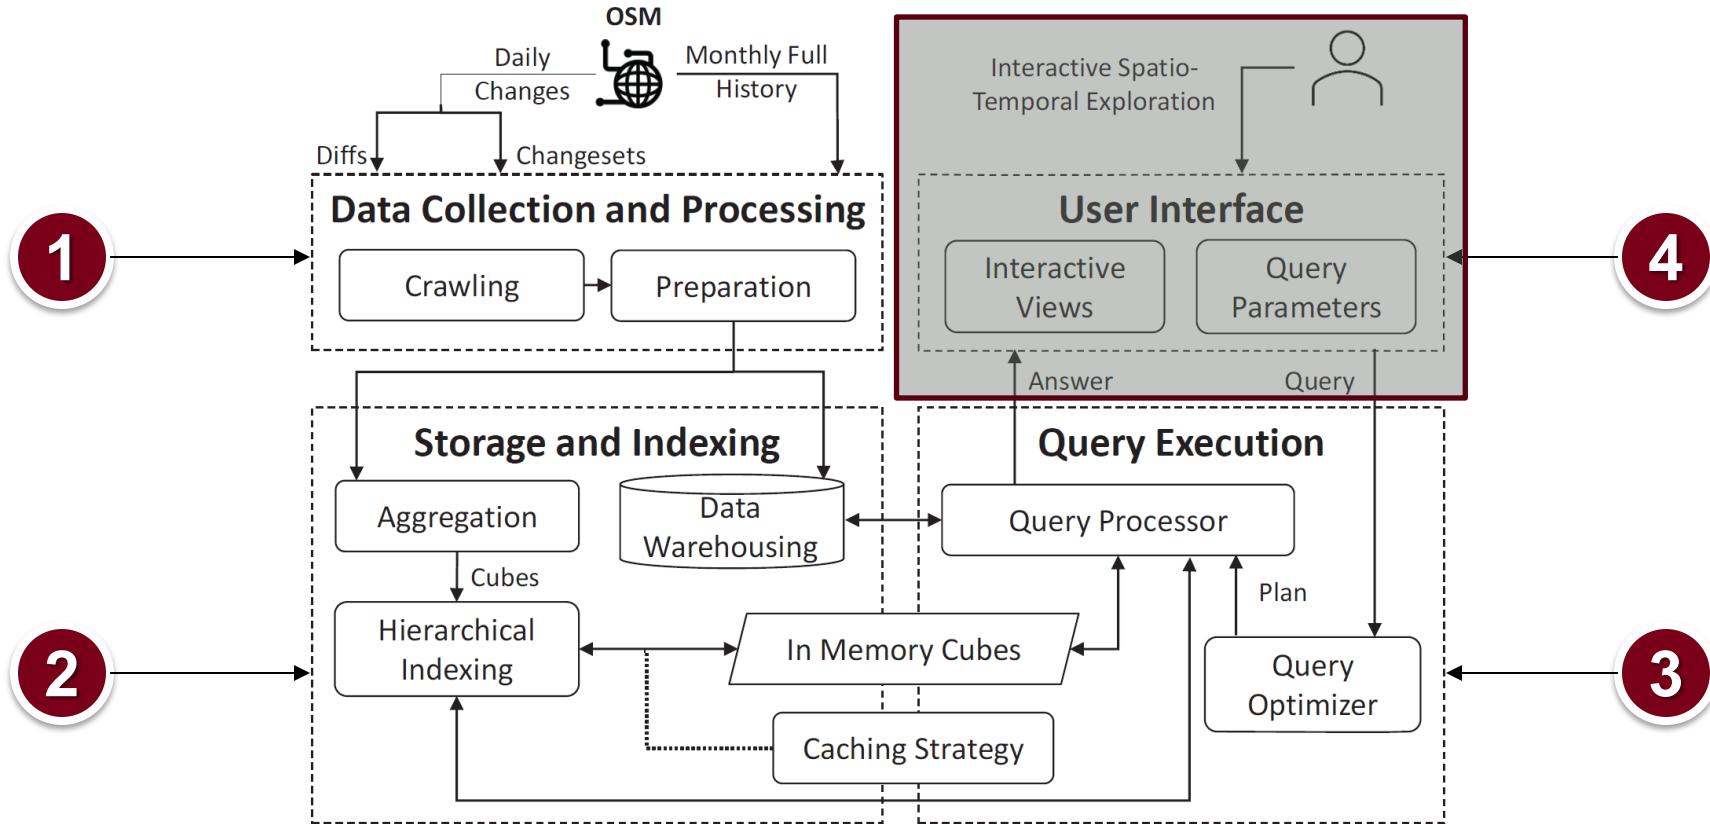
cost = 10 cubes

- A- All daily cost = 0 disk I/O
- B- 6 weekly + 4 daily cost = 6 disk I/O
- C- 1 monthly + 1 weekly + 8 daily cost = 2 disk I/O

# RASED Architecture



# RASED Architecture



# 3- User Interface

# 3- User Interface

- **Input:** Query Parameters

Start date  
2022-05-05

End date  
2022-06-04

Data is currently available from 2018-09-01 to 2022-06-04. Last 30 days are selected by default

**Category Selection**

Road Network

Navigational Metadata

POI Road Network

Bus

Cyclist

Pedestrian

Other

**Elements:**

Way

Relation

Node

**Operations:**

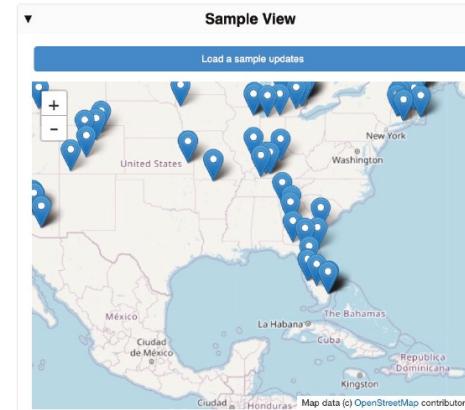
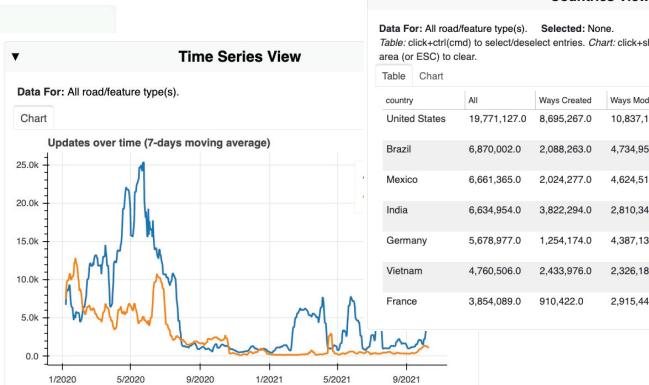
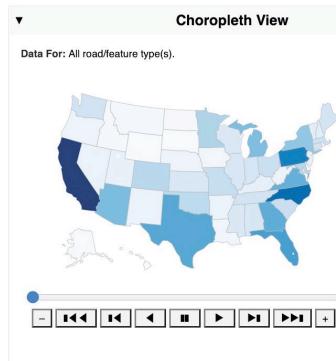
Create

Modify

**Query Data**

# 3- User Interface

- **Input:** Query Parameters
- **Output:**
  - Statistics either absolute numbers or percentages
  - Visualizations as:
    - Tabular format sorted on any column,
    - Various charts (bar, choropleth, time series)
    - Time-lapse video showing the road network evolution



Start date  
2022-05-05

End date  
2022-06-04

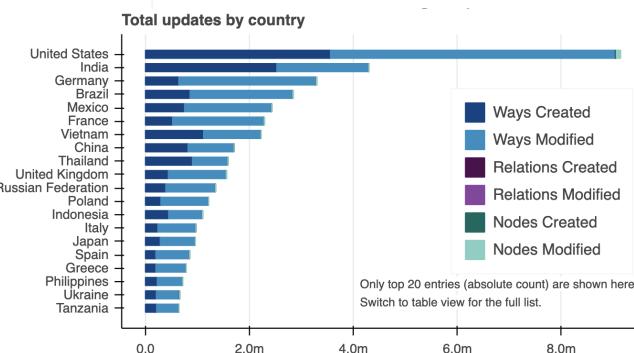
Data is currently available from 2018-09-01 to 2022-06-04. Last 30 days are selected by default

#### Category Selection

- Road Network
- Navigational Metadata
- POI Road Network
- Bus
- Cyclist
- Pedestrian
- Other

#### Elements:

- Way
- Relation
- Node



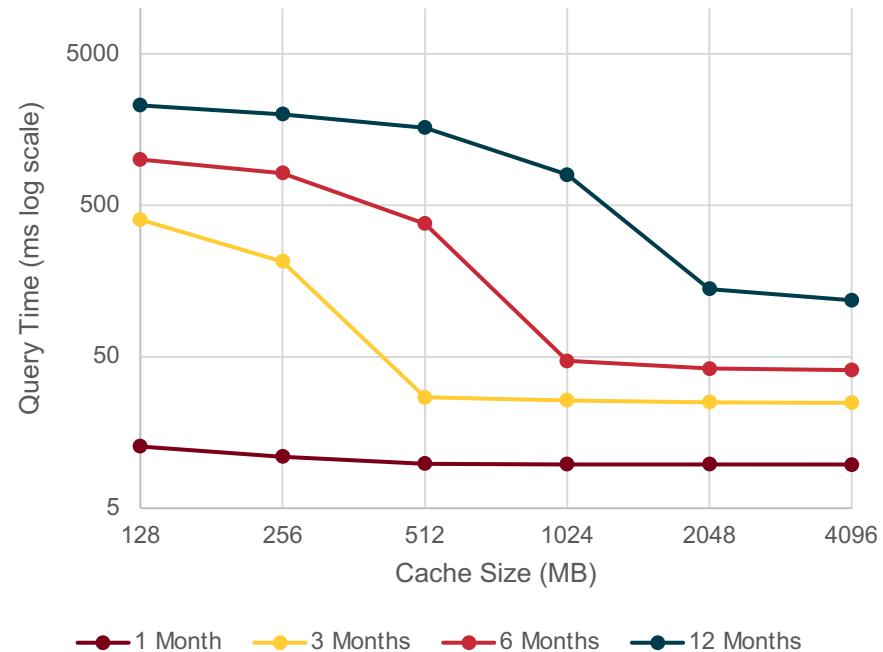
# Outline

- 1. Motivation**
- 2. Background**
  - Map updates, Datasets, and Queries
- 3. RASED Architecture & Main Modules**
  - Data Collection & Processing
  - Storage & Indexing
  - Query Execution
  - User Interface
- 4. Experimental Results**
- 5. RASED in Action (quick demo)**

# Setting RASED Parameters

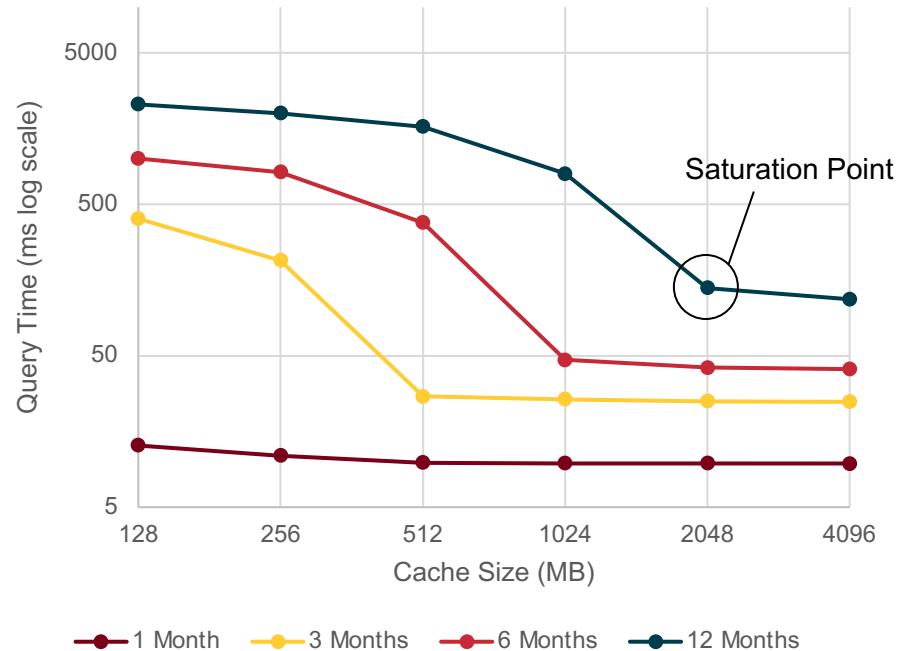
# Setting RASED Parameters

Setting RASED Cache Size



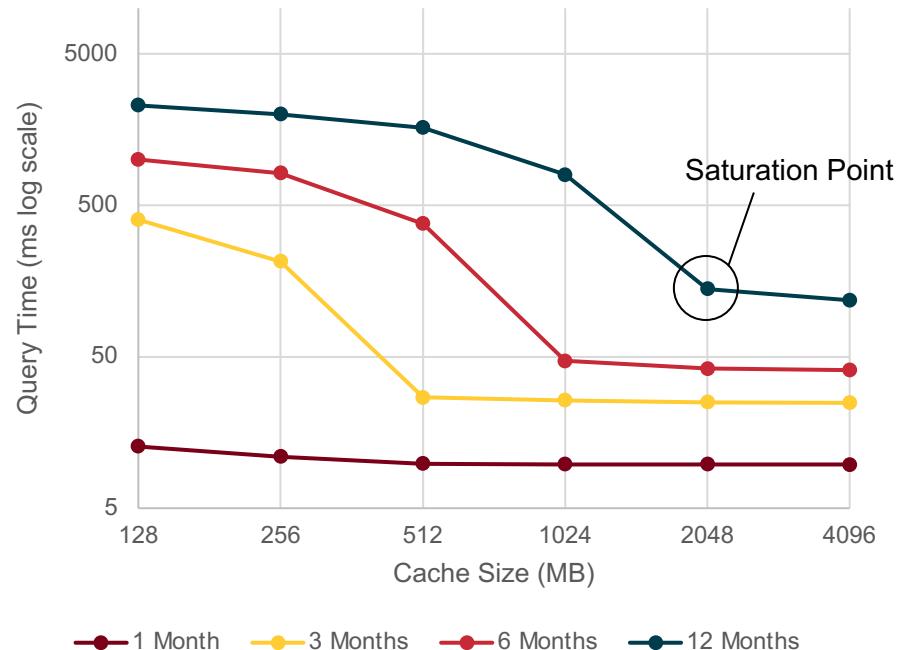
# Setting RASED Parameters

Setting RASED Cache Size



# Setting RASED Parameters

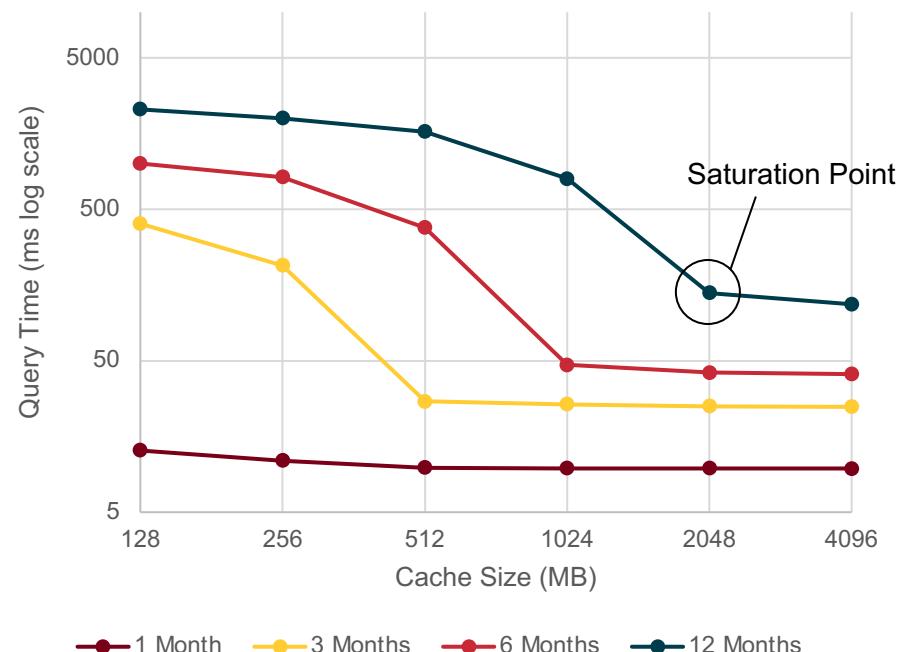
Setting RASED Cache Size



RASED Uses 2048 MB for Cache

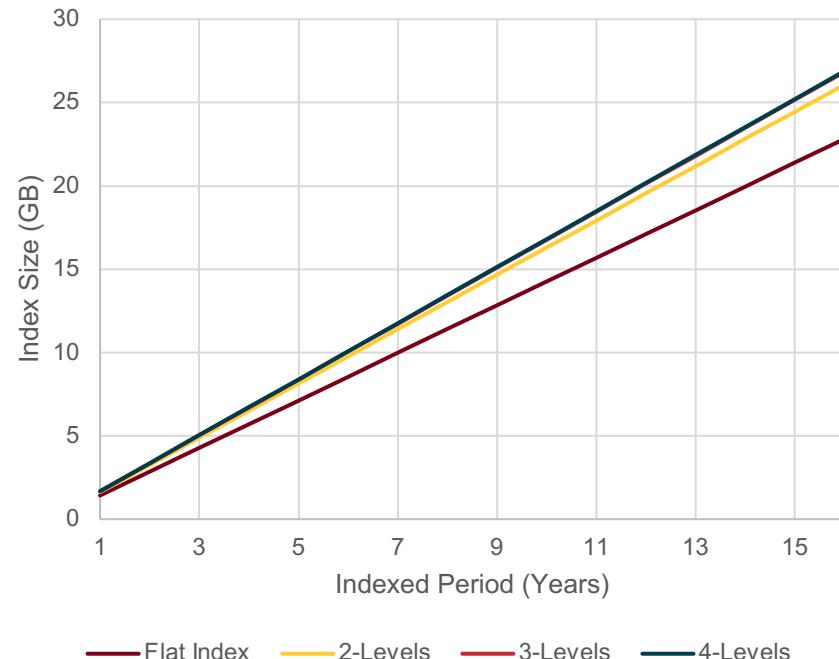
# Setting RASED Parameters

Setting RASED Cache Size



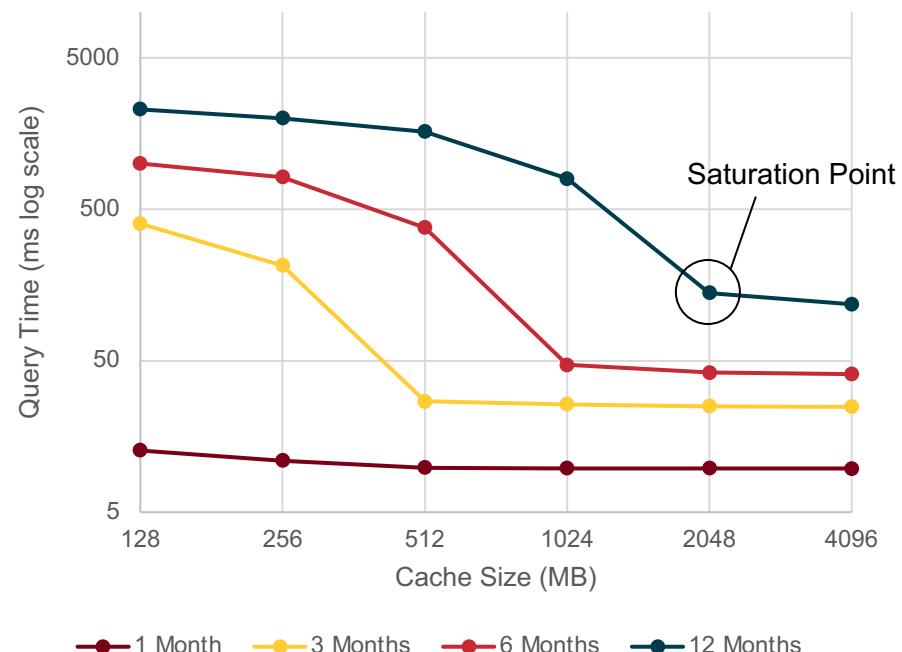
RASED Uses 2048 MB for Cache

Setting RASED Number of Levels



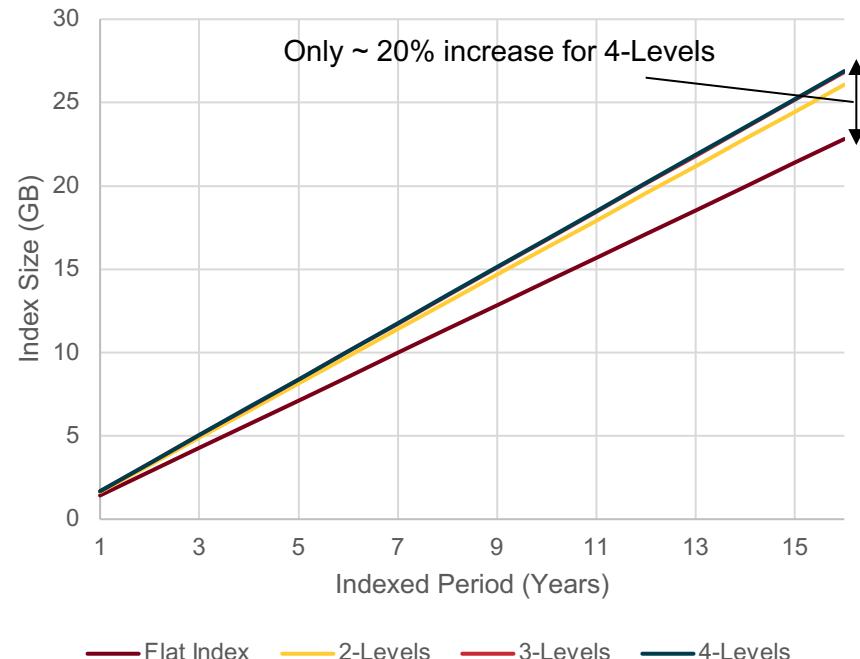
# Setting RASED Parameters

Setting RASED Cache Size



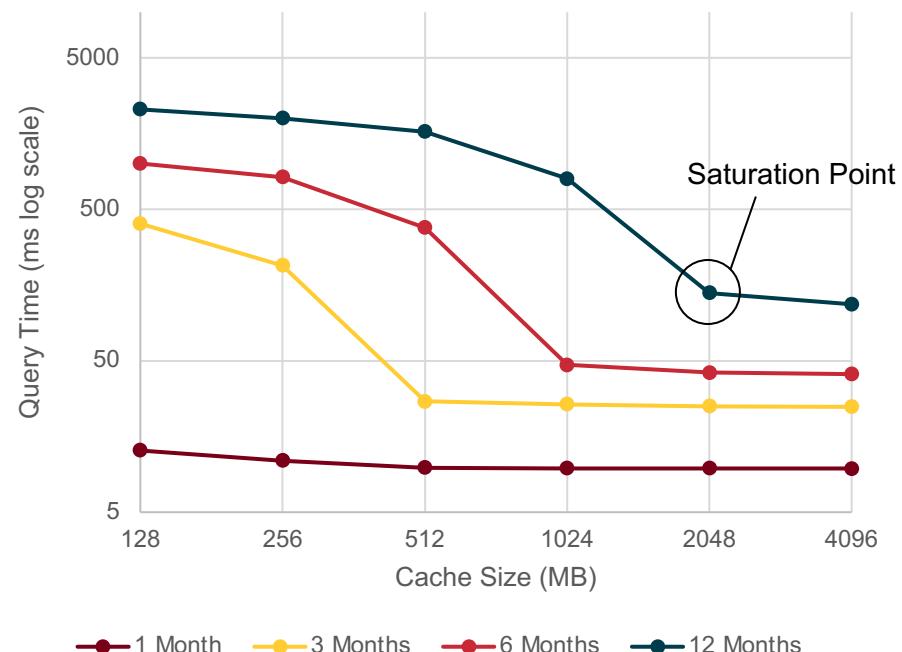
RASED Uses 2048 MB for Cache

Setting RASED Number of Levels



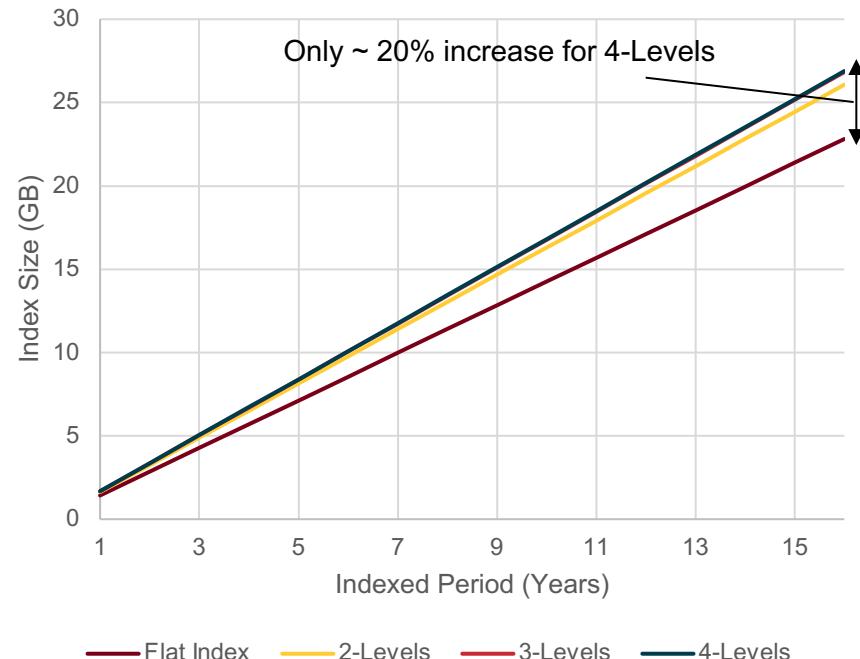
# Setting RASED Parameters

Setting RASED Cache Size



RASED Uses 2048 MB for Cache

Setting RASED Number of Levels



RASED Index Uses 4 Levels

# RASED Performance

# RASED Performance

**RASED-F:** one level flat index, neither caching nor level optimization

**RASED-O:** hierachal index with level optimization, no caching

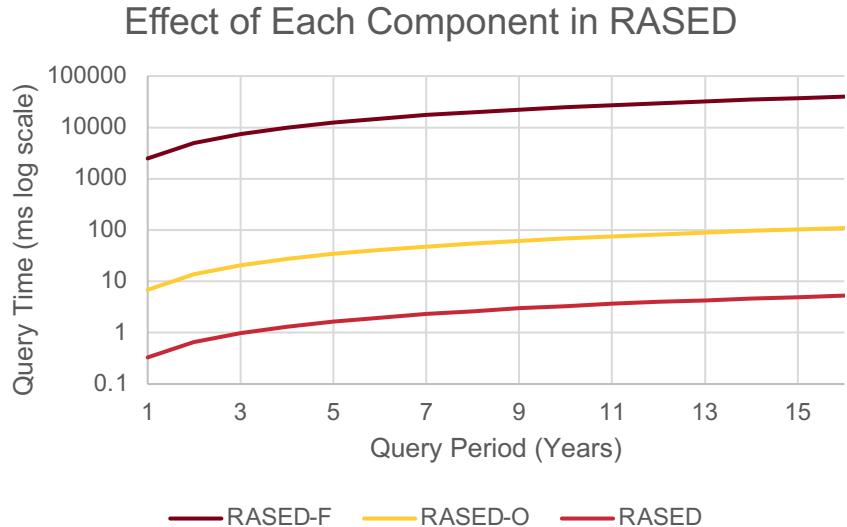
**RASED:** full RASED system with both caching nor level optimization

# RASED Performance

**RASED-F:** one level flat index, neither caching nor level optimization

**RASED-O:** hierachal index with level optimization, no caching

**RASED:** full RASED system with both caching nor level optimization

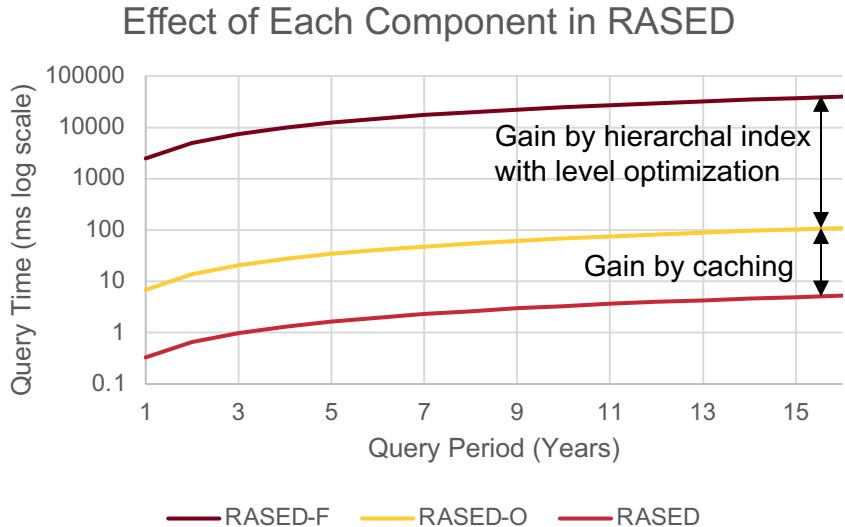


# RASED Performance

**RASED-F:** one level flat index, neither caching nor level optimization

**RASED-O:** hierachal index with level optimization, no caching

**RASED:** full RASED system with both caching and level optimization

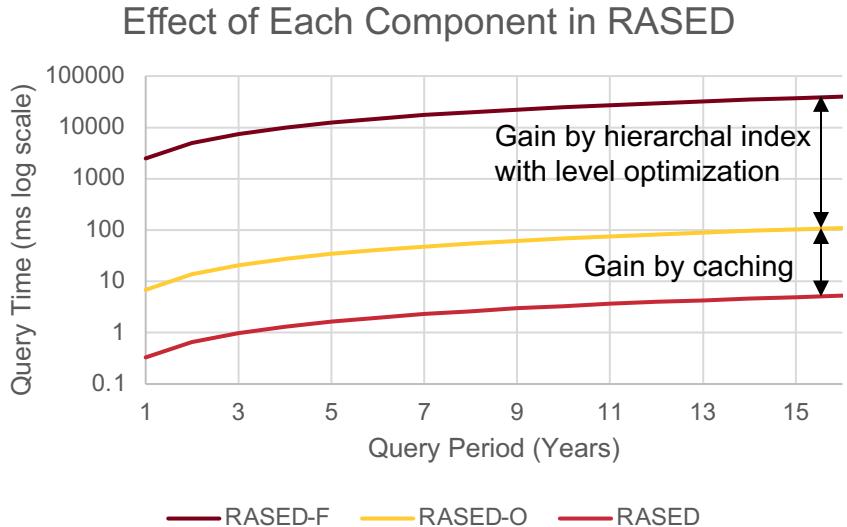


# RASED Performance

**RASED-F:** one level flat index, neither caching nor level optimization

**RASED-O:** hierachal index with level optimization, no caching

**RASED:** full RASED system with both caching and level optimization



RASED indexing, optimizing, and caching reduce response times by orders of magnitude

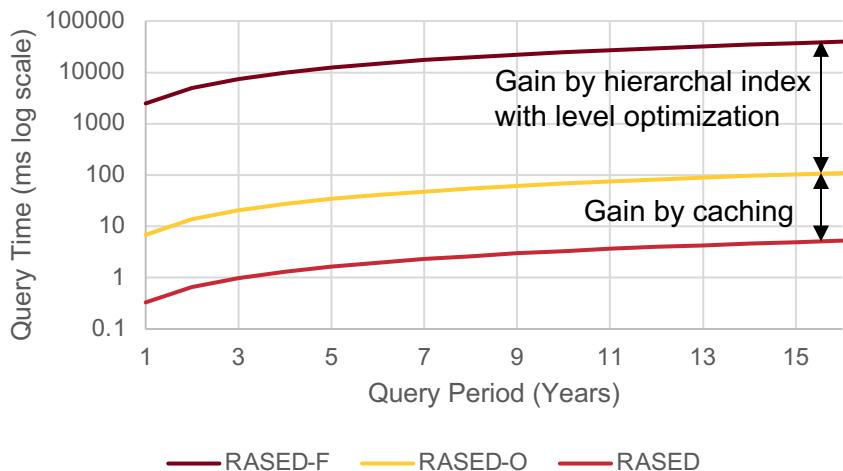
# RASED Performance

**RASED-F:** one level flat index, neither caching nor level optimization

**RASED-O:** hierachal index with level optimization, no caching

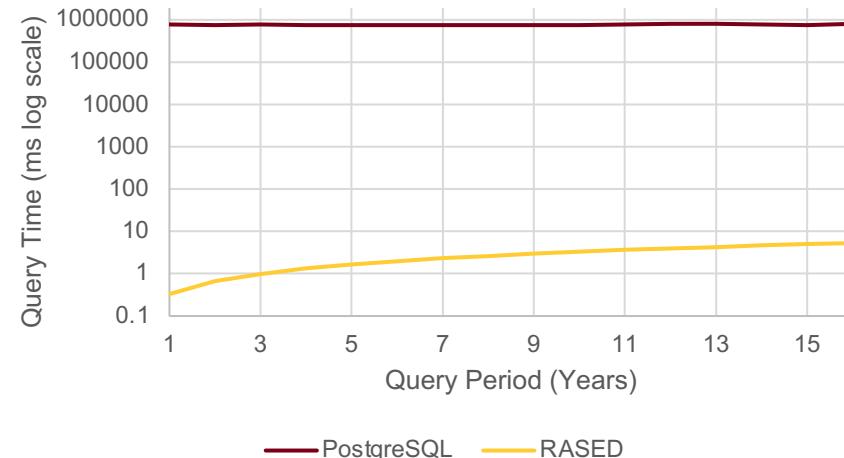
**RASED:** full RASED system with both caching and level optimization

Effect of Each Component in RASED



**PostgreSQL** constantly requires scanning the whole data because query involves multiple attributes in the Group By clause

Overall Performance



RASED indexing, optimizing, and caching reduce response times by orders of magnitude

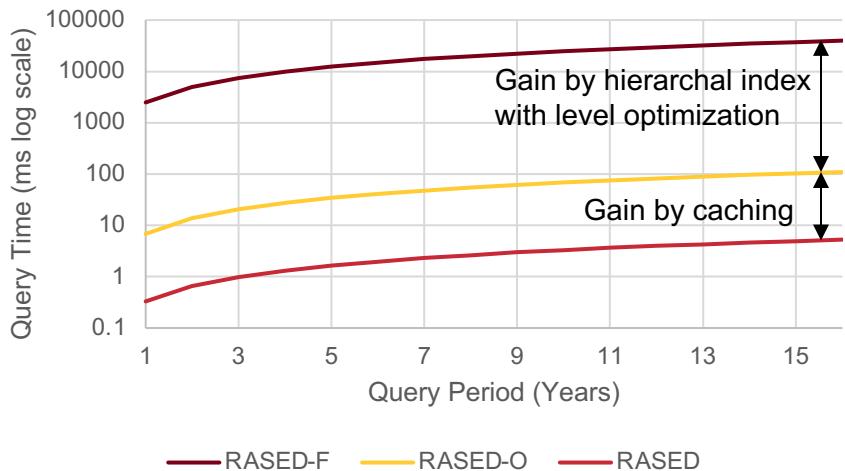
# RASED Performance

**RASED-F:** one level flat index, neither caching nor level optimization

**RASED-O:** hierachal index with level optimization, no caching

**RASED:** full RASED system with both caching and level optimization

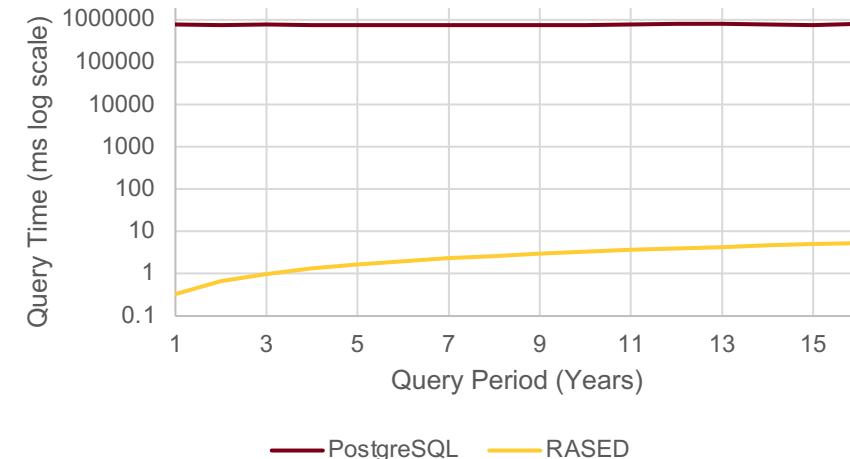
Effect of Each Component in RASED



RASED indexing, optimizing, and caching reduce response times by orders of magnitude

**PostgreSQL** constantly requires scanning the whole data because query involves multiple attributes in the Group By clause

Overall Performance



RASED queries are always answered in milliseconds

# Outline

## 1. Motivation

## 2. Background

- Map updates, Datasets, and Queries

## 3. RASED Architecture & Main Modules

- Data Collection & Processing
- Storage & Indexing
- Query Execution
- User Interface

## 4. Experimental Results

## 5. RASED in Action (quick demo)

# DEMO

<https://rased.cs.umn.edu>

[https://www.youtube.com/  
watch?v=fB2JxTZ60wU](https://www.youtube.com/watch?v=fB2JxTZ60wU)

THANK YOU